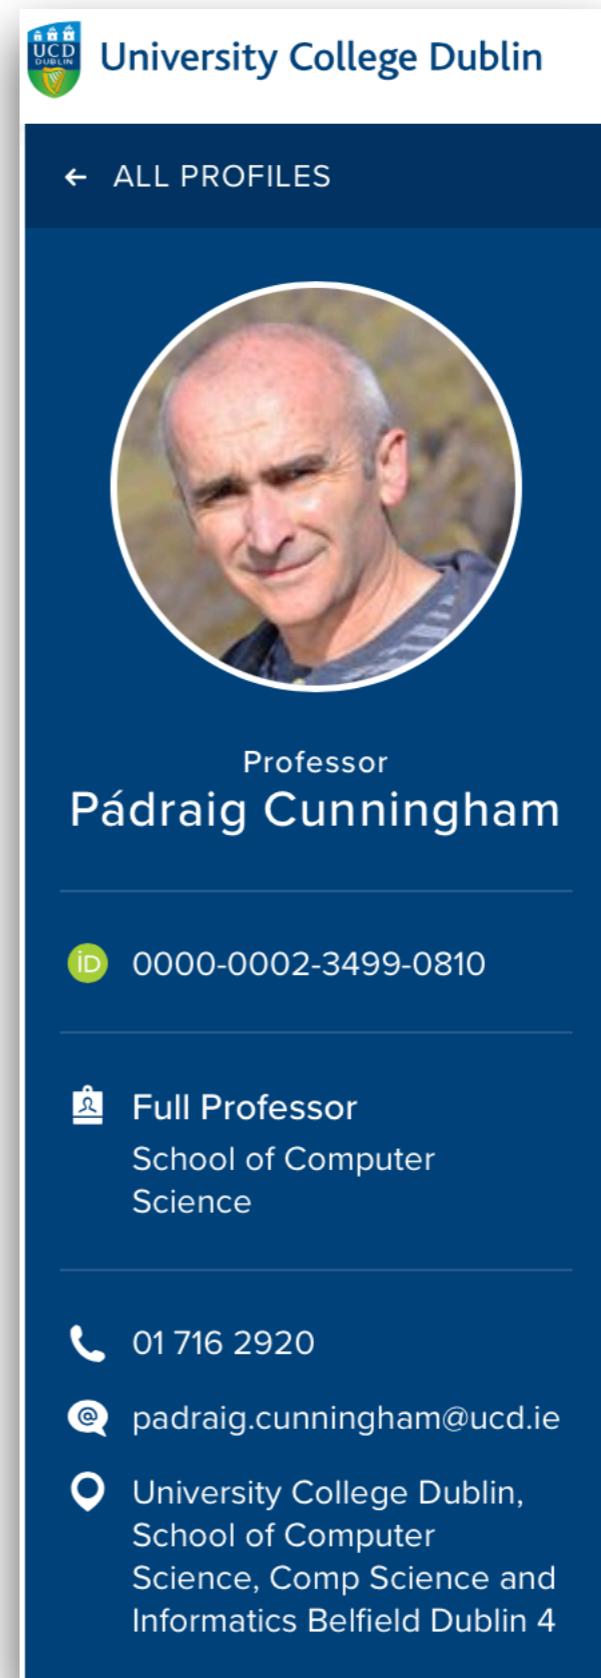
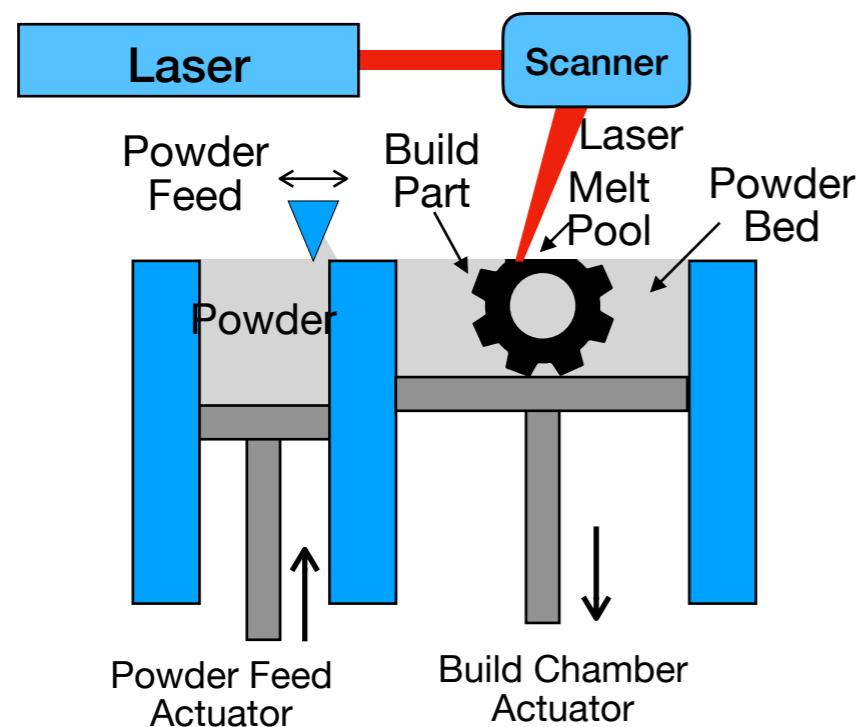


Time-Series Classification

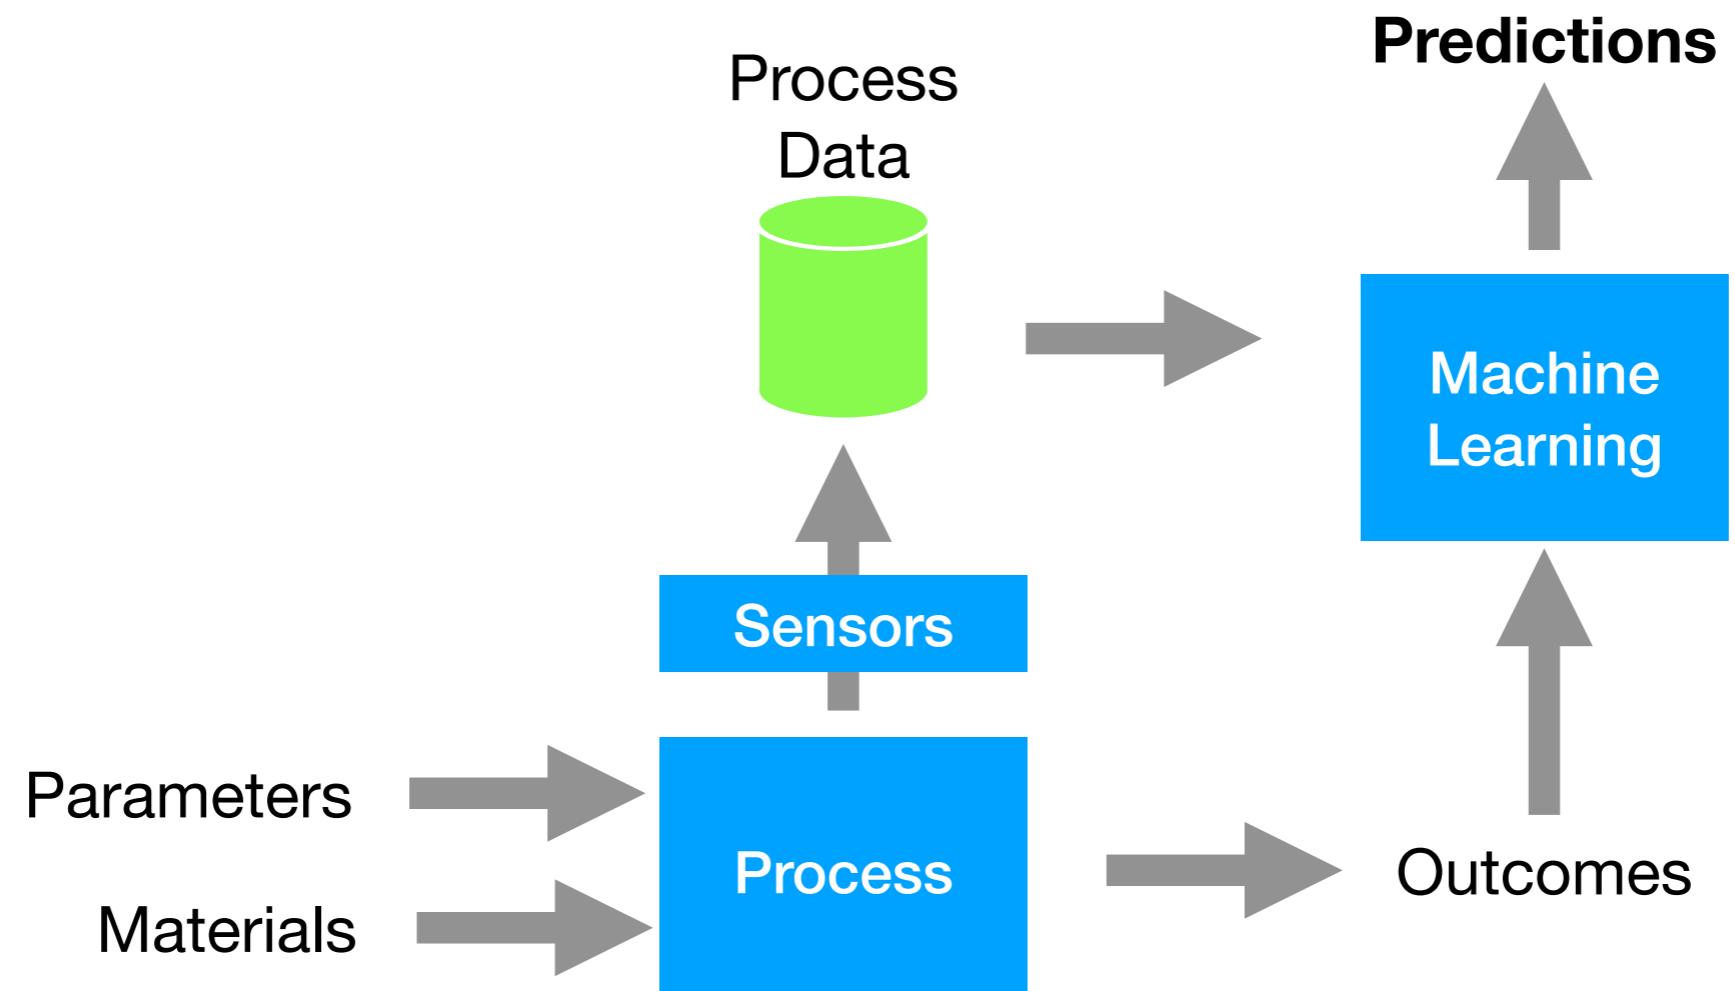
- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection



Metal 3D Printing



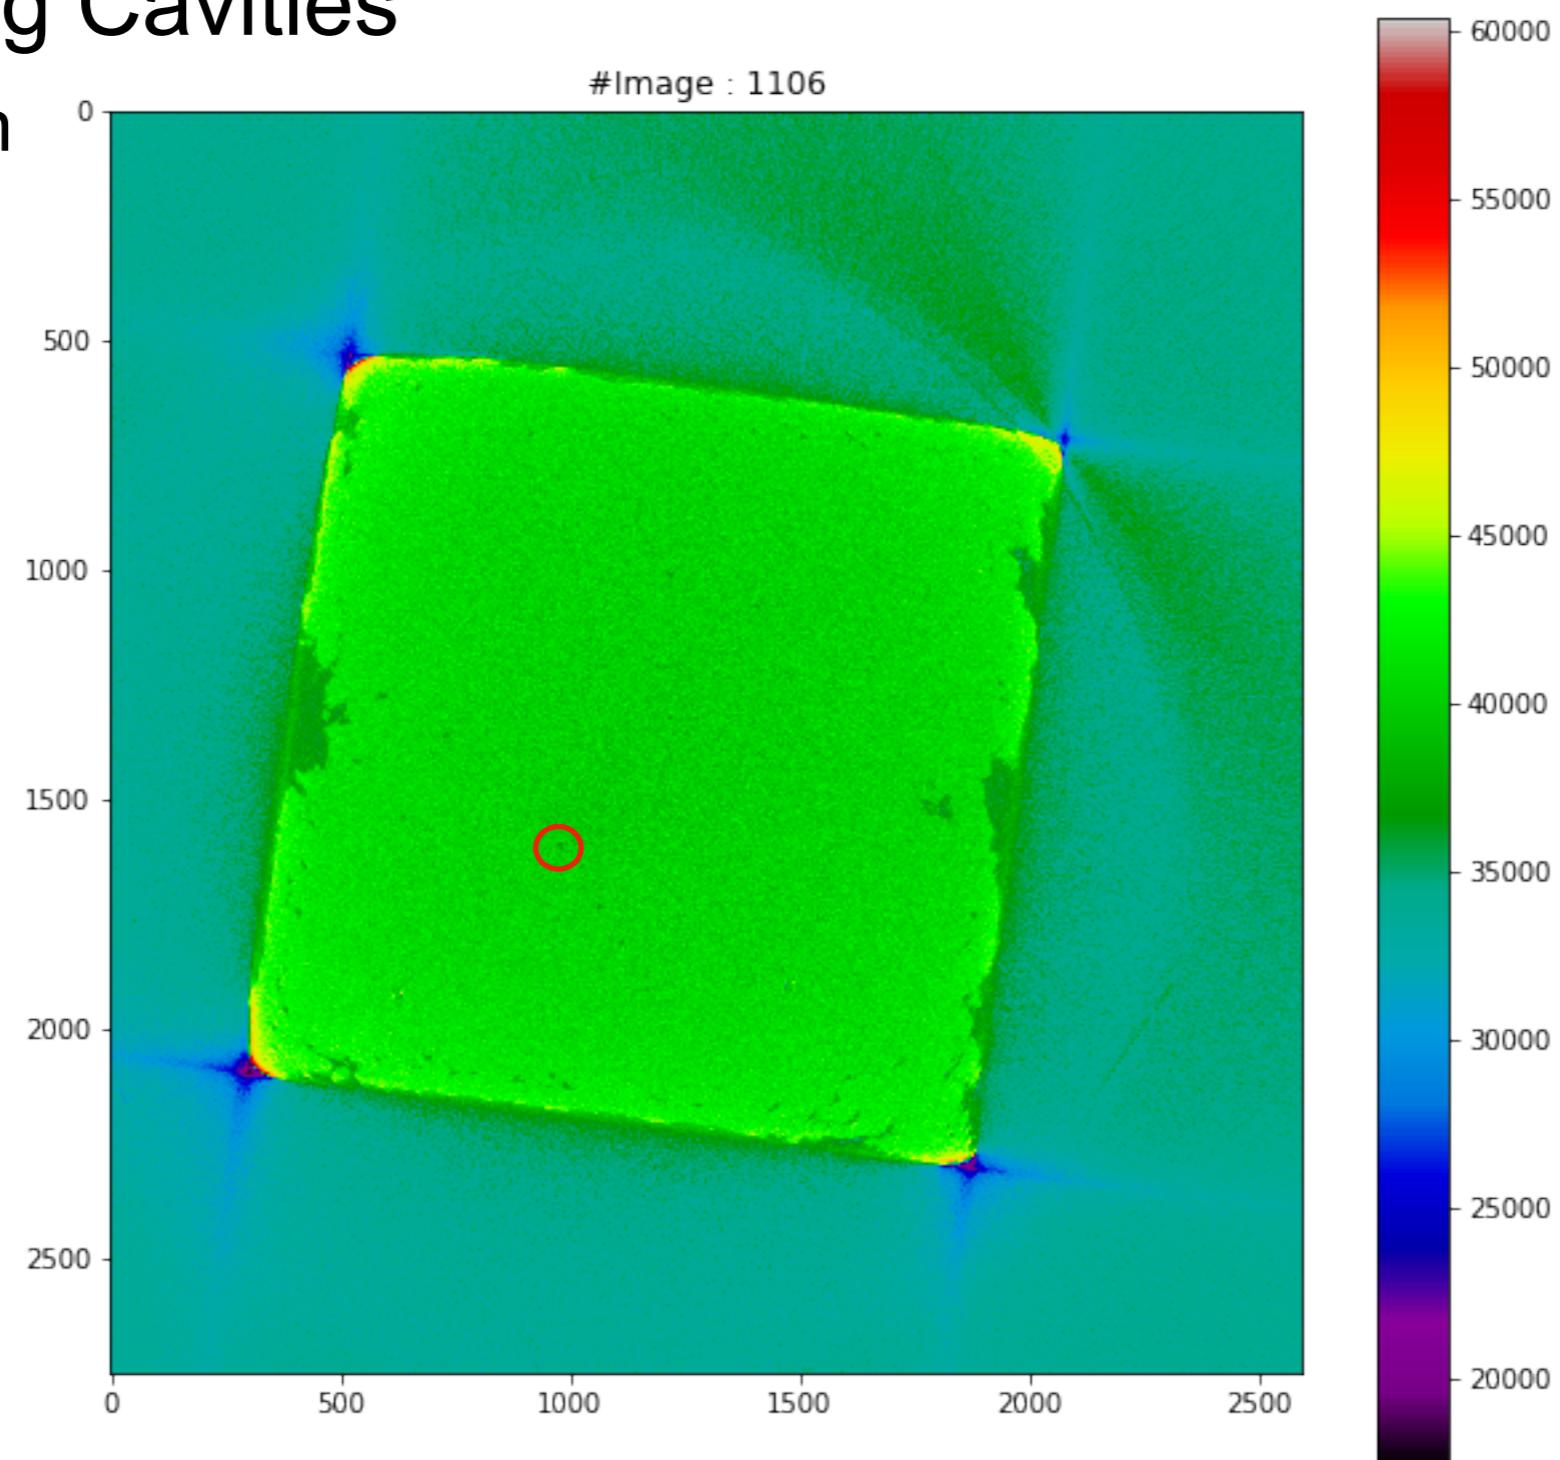
ML & Process Monitoring



Example: 3D Printing

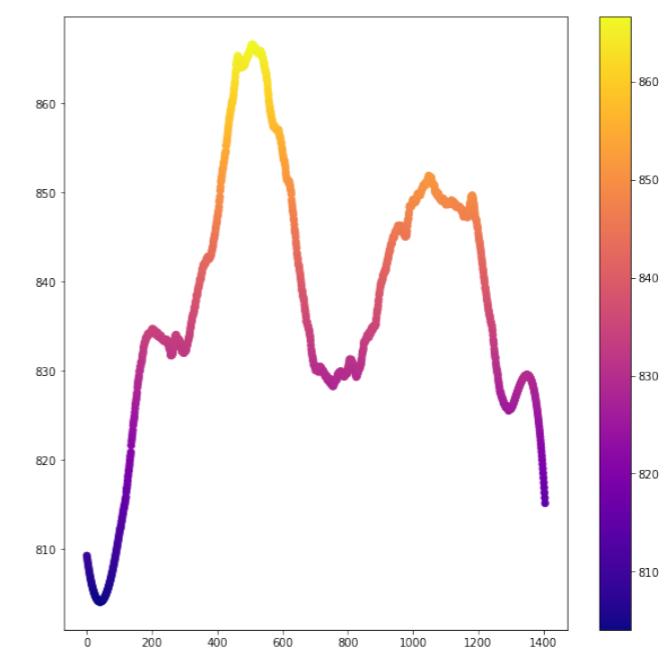
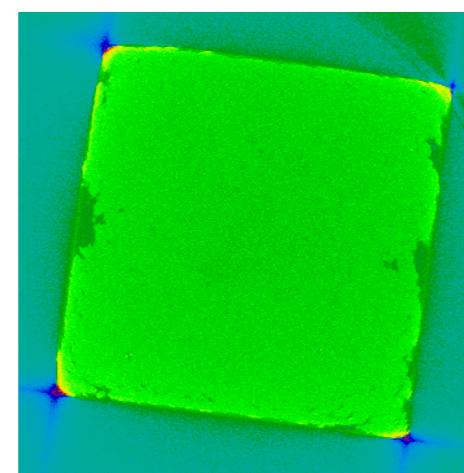
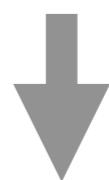
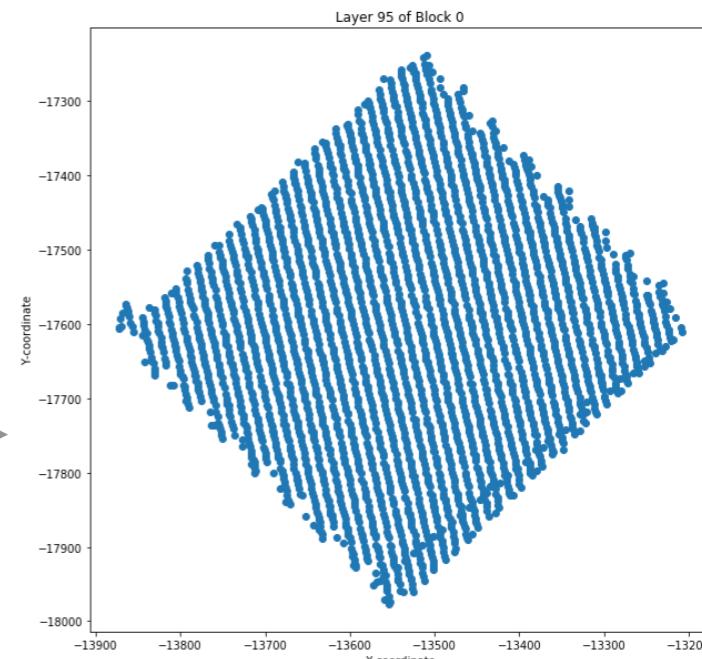
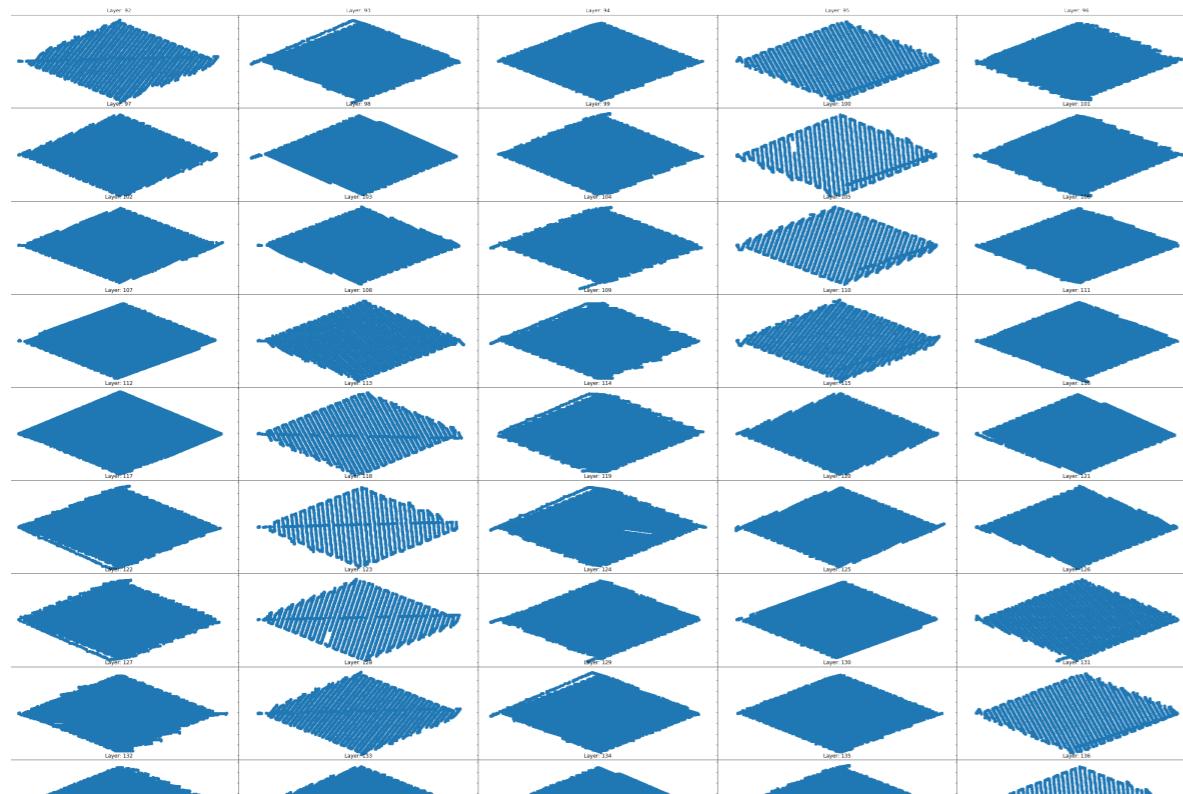
■ Predicting Cavities

CT Scan



Example: 3D printing

- Learn temperature profiles associated with anomalies

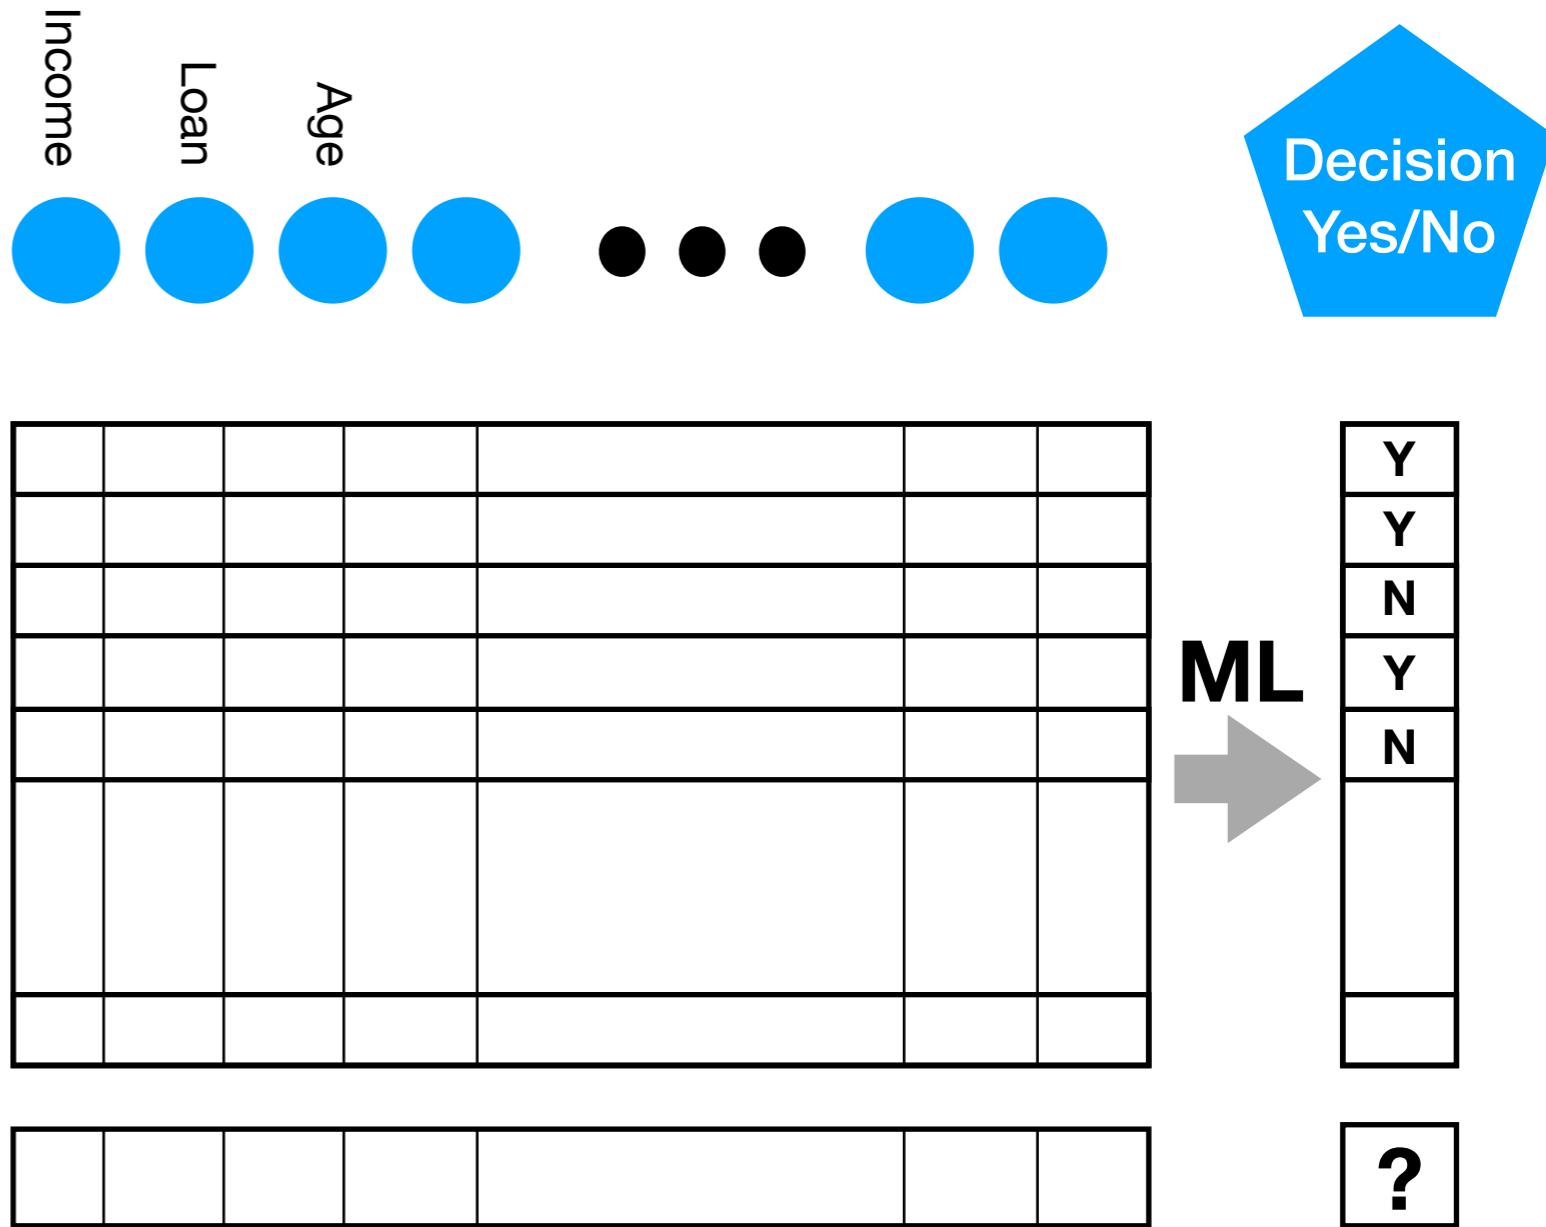


Time-Series Classification

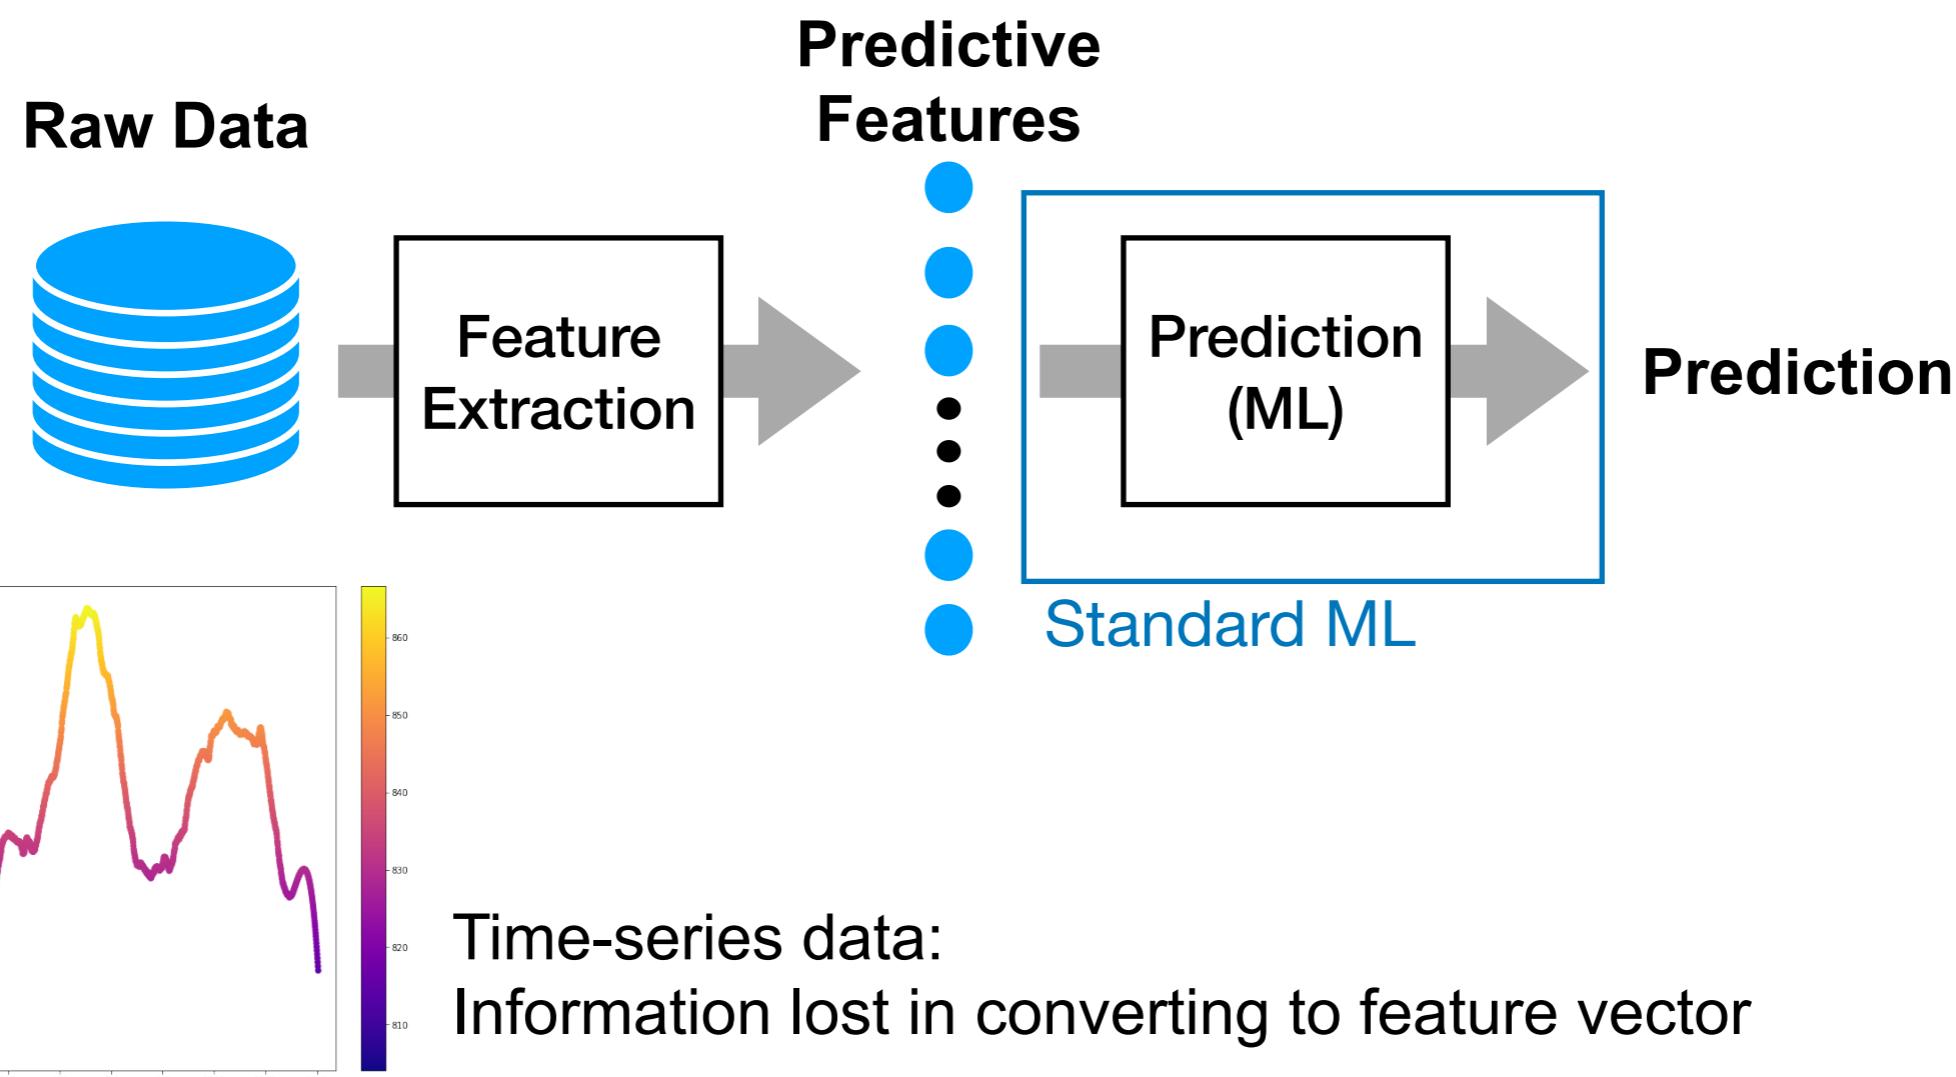
- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection

Learning from Data

- Table of historic data
- Each row
 - description of case
 - decision (Yes/No)



Supervised ML (aka Predictive Analytics)



Time-Series Classification

- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection

Eager v Lazy Classifiers

- **Eager Learning Classification Strategy**
 - Classifier builds a full model during an initial training phase, to use later when new query examples arrive.
 - More offline setup work, less work at run-time.
 - Generalise before seeing the query example.
- **Lazy Learning Classification Strategy**
 - Classifier keeps all the training examples for later use.
 - Little work is done offline, wait for new query examples.
 - Focus on the local space around the examples.
- **Distance-based Models:** Many learning algorithms are based on generalising from training data to unseen data by exploiting the distances (or similarities) between the two.

Example: Athlete Selection

- **Training set** of performance ratings for 20 college athletes, where each athlete is described by 2 continuous features: *speed*, *agility*.
- Each athlete has a **target class label** indicating whether they were selected for the university athletics team: 'Yes' or 'No'.

Athlete	Speed	Agility	Selected
x1	2.50	6.00	No
x2	3.75	8.00	No
x3	2.25	5.50	No
x4	3.25	8.25	No
x5	2.75	7.50	No
x6	4.50	5.00	No
x7	3.50	5.25	No
x8	3.00	3.25	No
x9	4.00	4.00	No
x10	4.25	3.75	No

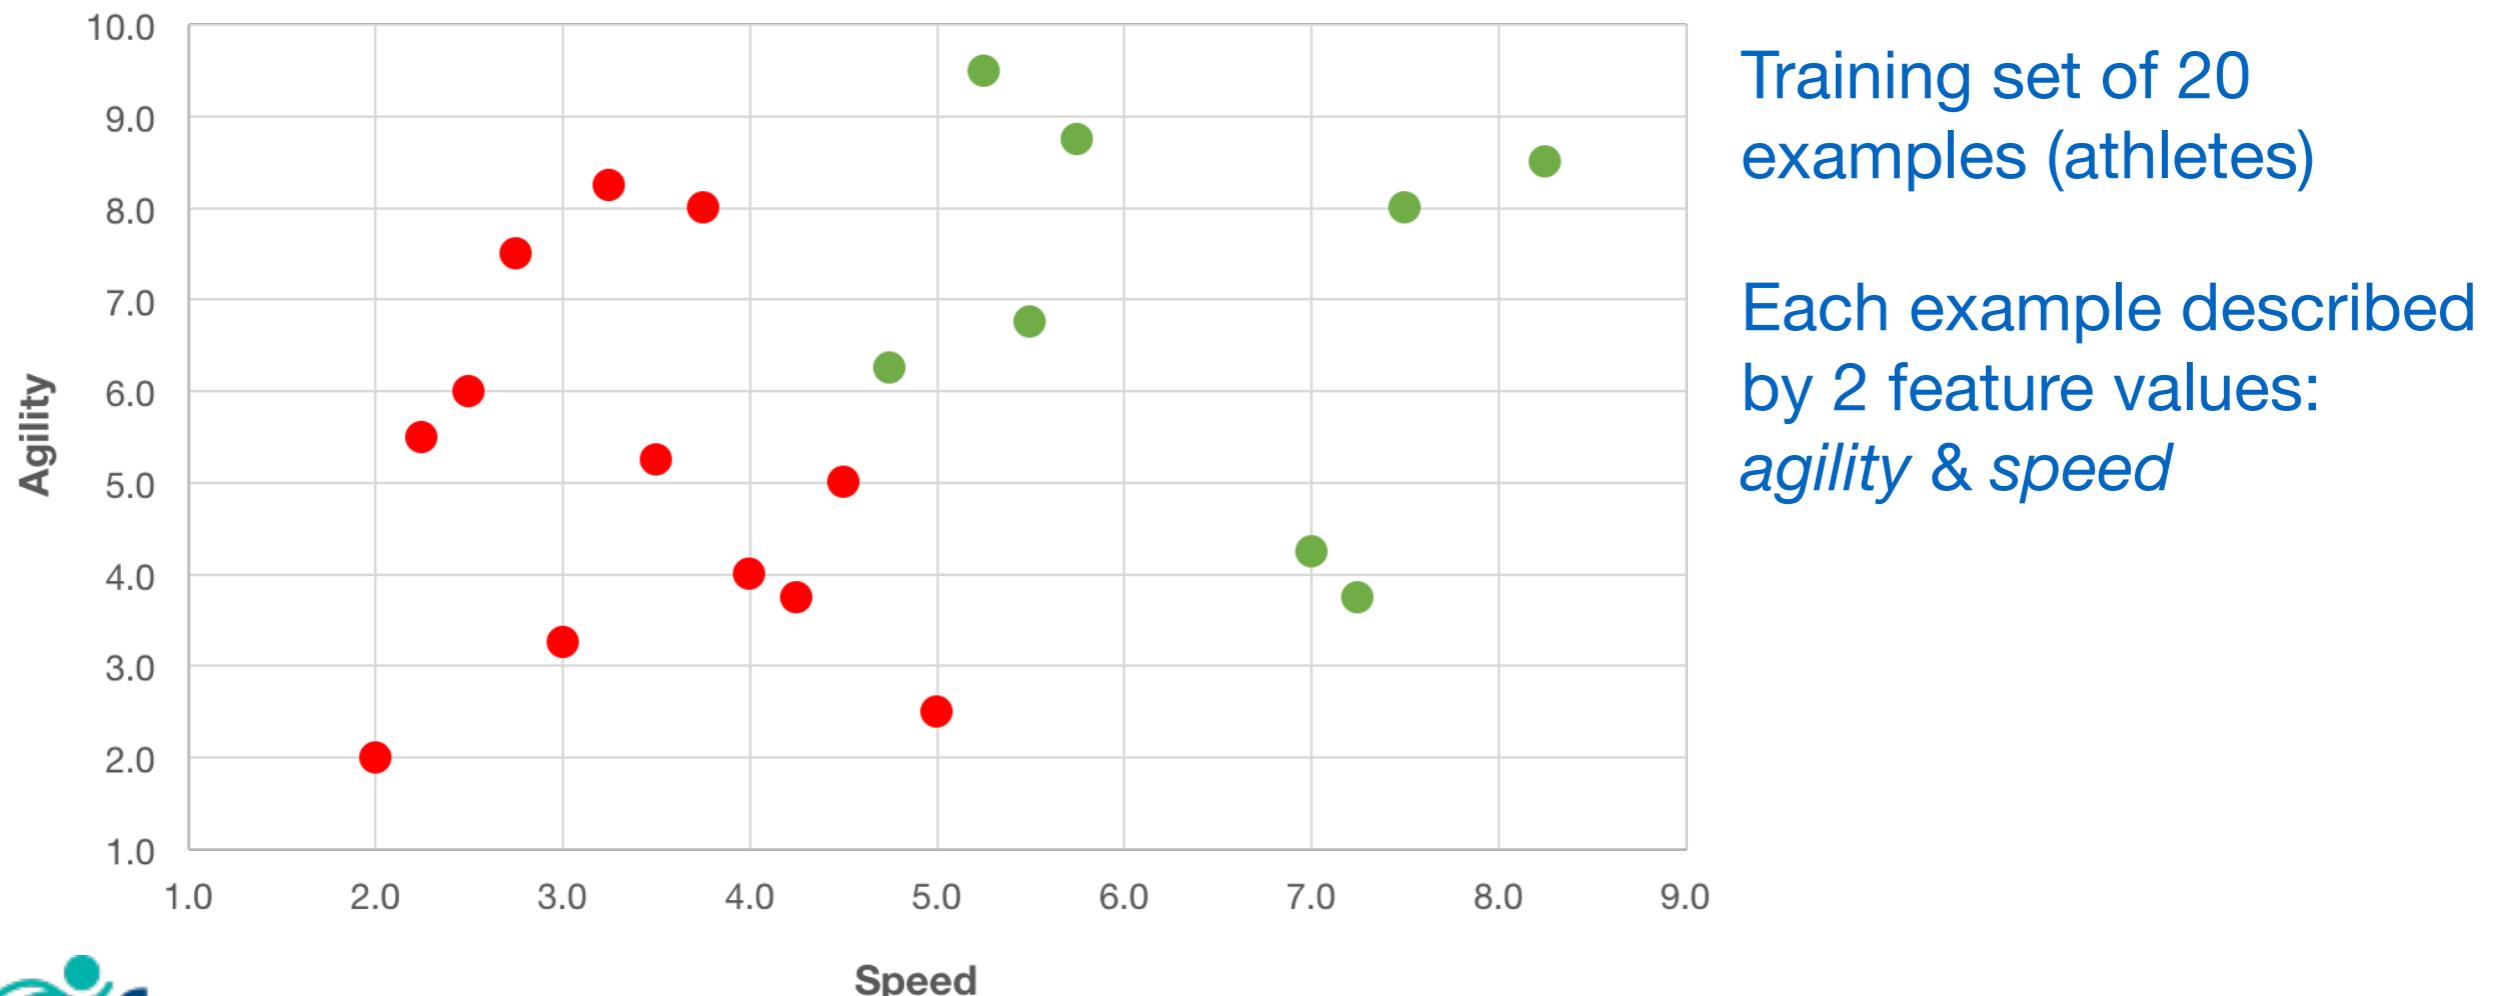
Athlete	Speed	Agility	Selected
x11	2.00	2.00	No
x12	5.00	2.50	No
x13	8.25	8.50	Yes
x14	5.75	8.75	Yes
x15	4.75	6.25	Yes
x16	5.50	6.75	Yes
x17	5.25	9.50	Yes
x18	7.00	4.25	Yes
x19	7.50	8.00	Yes
x20	7.25	3.75	Yes

Q. Will a new athlete q be selected: 'Yes' or 'No'?

Athlete	Speed	Agility	Selected
q	3.00	8.00	???

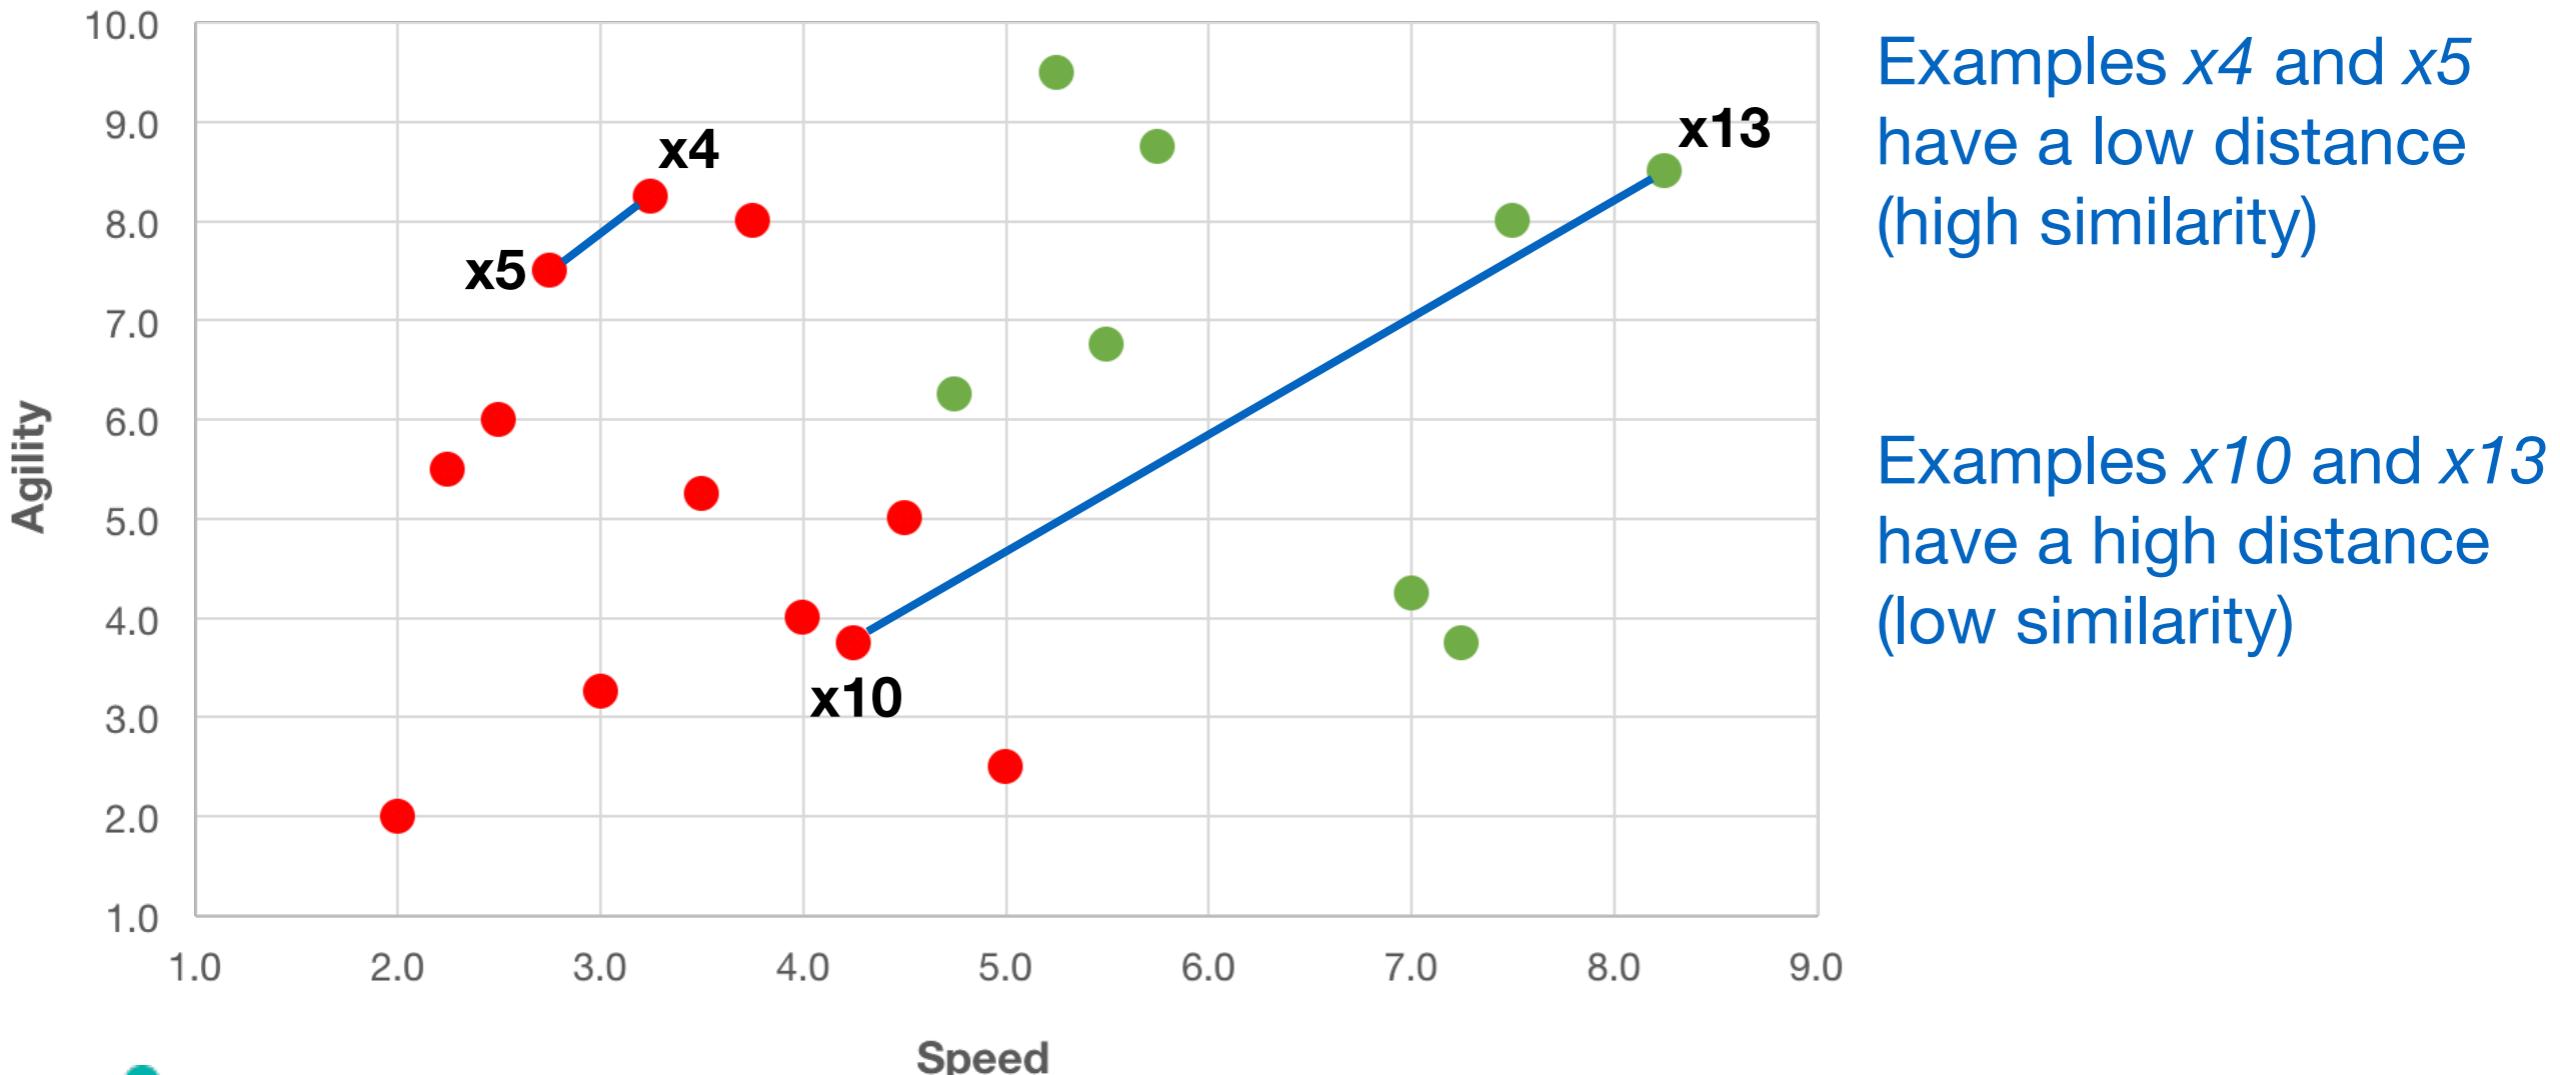
Feature Spaces

- A **feature space** is a D -dimensional coordinate space used to represent the input examples for a given problem, with one coordinate for each descriptive feature.
- **Example:** Use a feature space to visually position the 20 athletes in a 2-dimensional coordinate space (i.e. *agility* versus *speed*):



Measuring Distance

- Measuring the **distance** (or **similarity**) between two examples is fundamental to many ML algorithms.
- Many measures can be used to calculate distance. There is no “best” distance measure. The choice is highly problem-dependent.



Measuring Distance

- **Euclidean distance:** Most common measure used to quantify distance between two examples with numeric features.
- Given by the "straight line" distance between two points in a Euclidean coordinate space - e.g. a feature space.
- Calculated as the square root of sum of squared differences for each feature f representing a pair of examples.
- The output is a real value ≥ 0 , where a larger value indicates two examples are more distant (i.e. less similar to one another).

Input:
2 examples
 \mathbf{p} and \mathbf{q}

$$ED(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{f \in F} (q_f - p_f)^2}$$

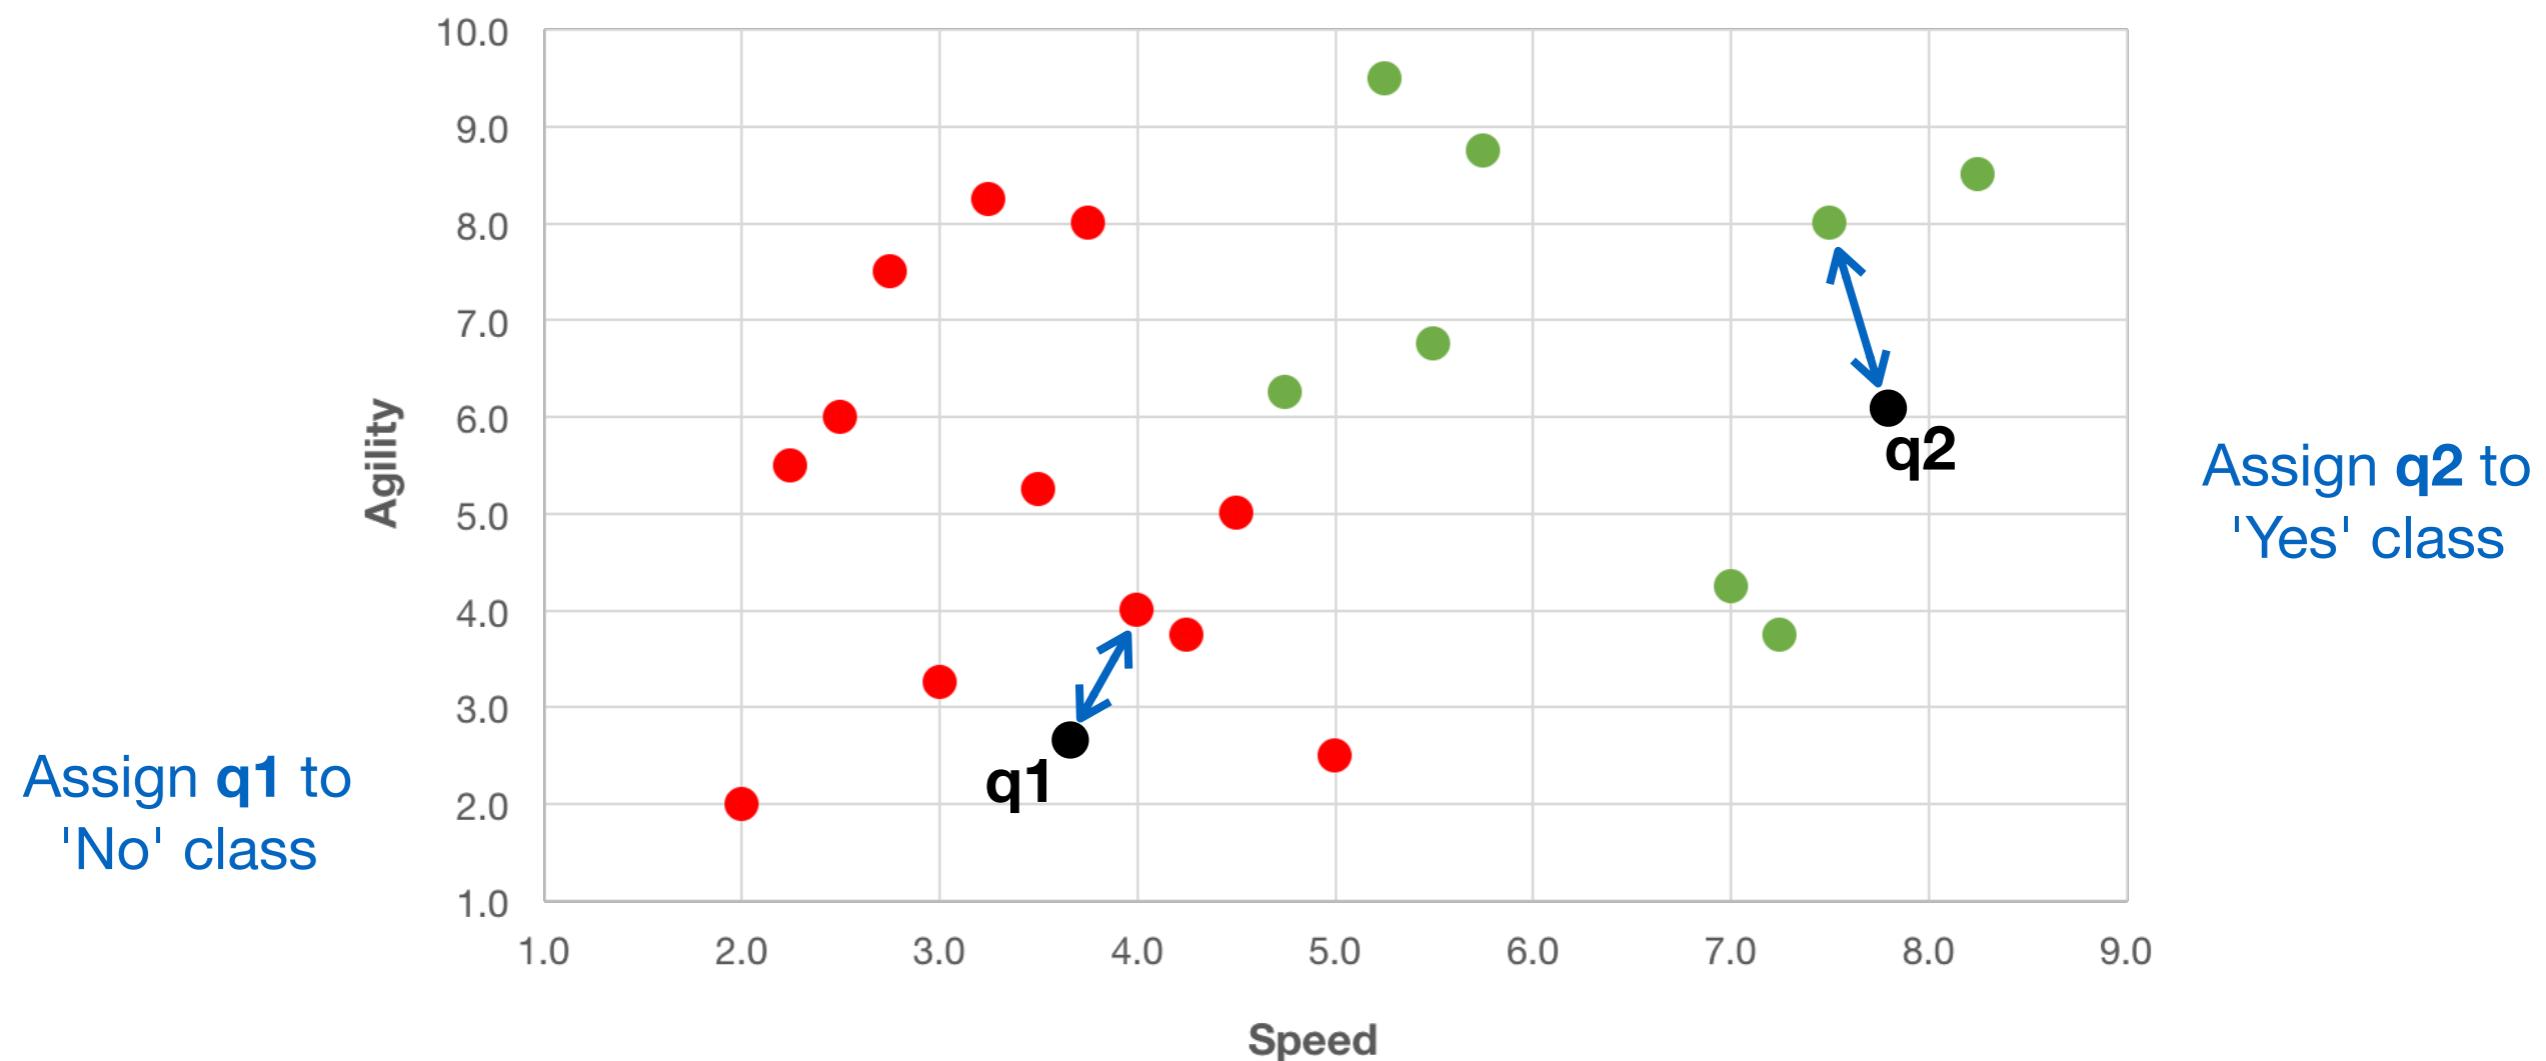
Calculate square of the difference between the examples on feature f

For each feature f in the full set of features F

Nearest Neighbour Classifier

Lazy learning approach: Do not build a model for the data. Identify most similar previous example(s) from the training set for which a label has already been assigned, using some distance function.

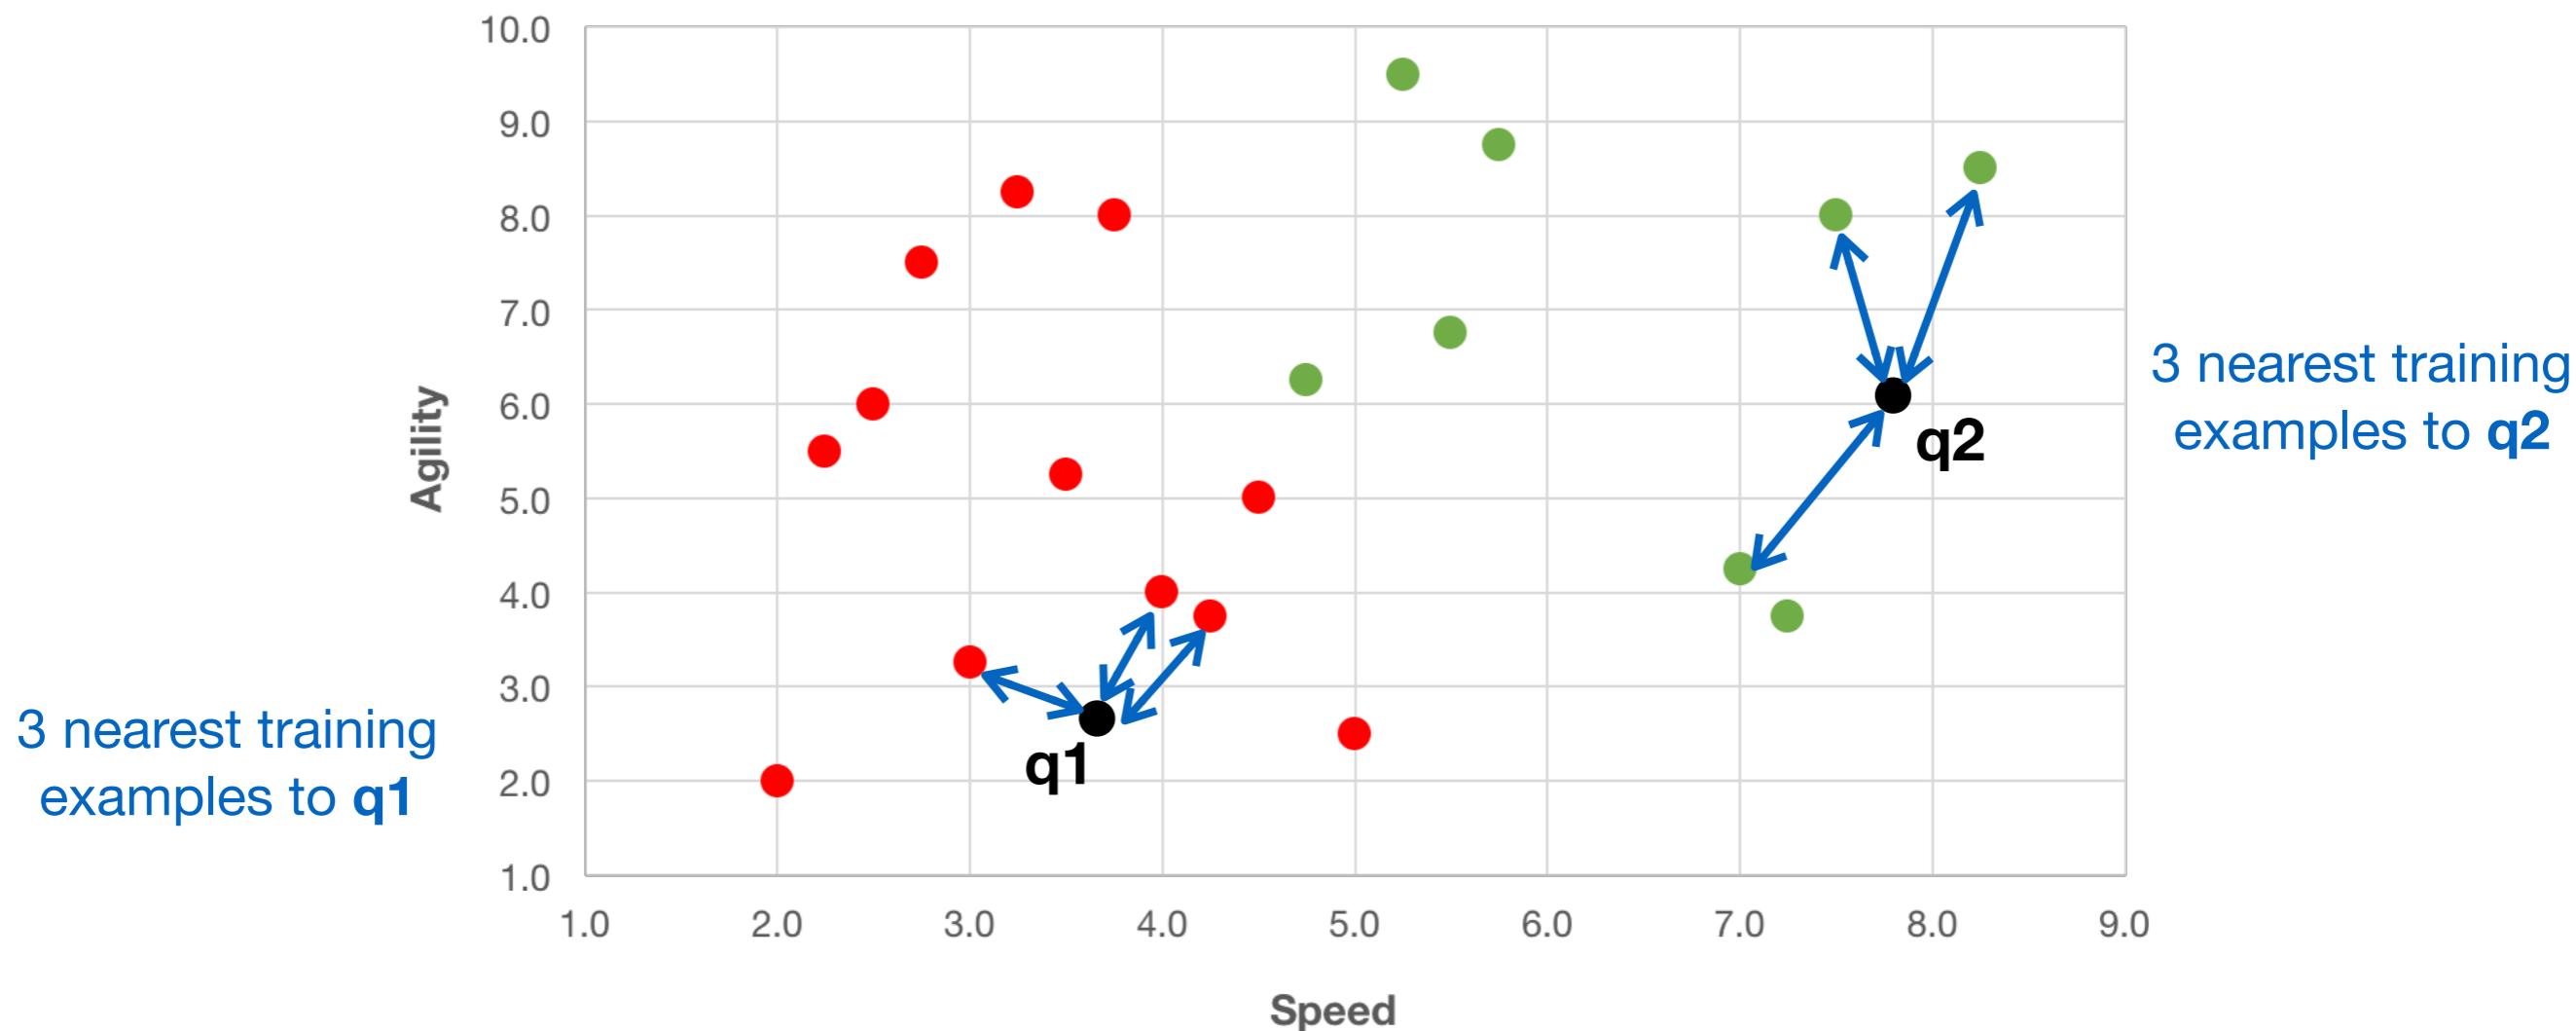
Nearest neighbour rule (1NN): For a new query input q , find a single labelled example x closest to q , and assign q the same label as x .



k -Nearest Neighbour Classifier

k -Nearest neighbours (kNN): The NN approach naturally generalises to the case where we use k nearest neighbours from the training set to assign a label to a new query input.

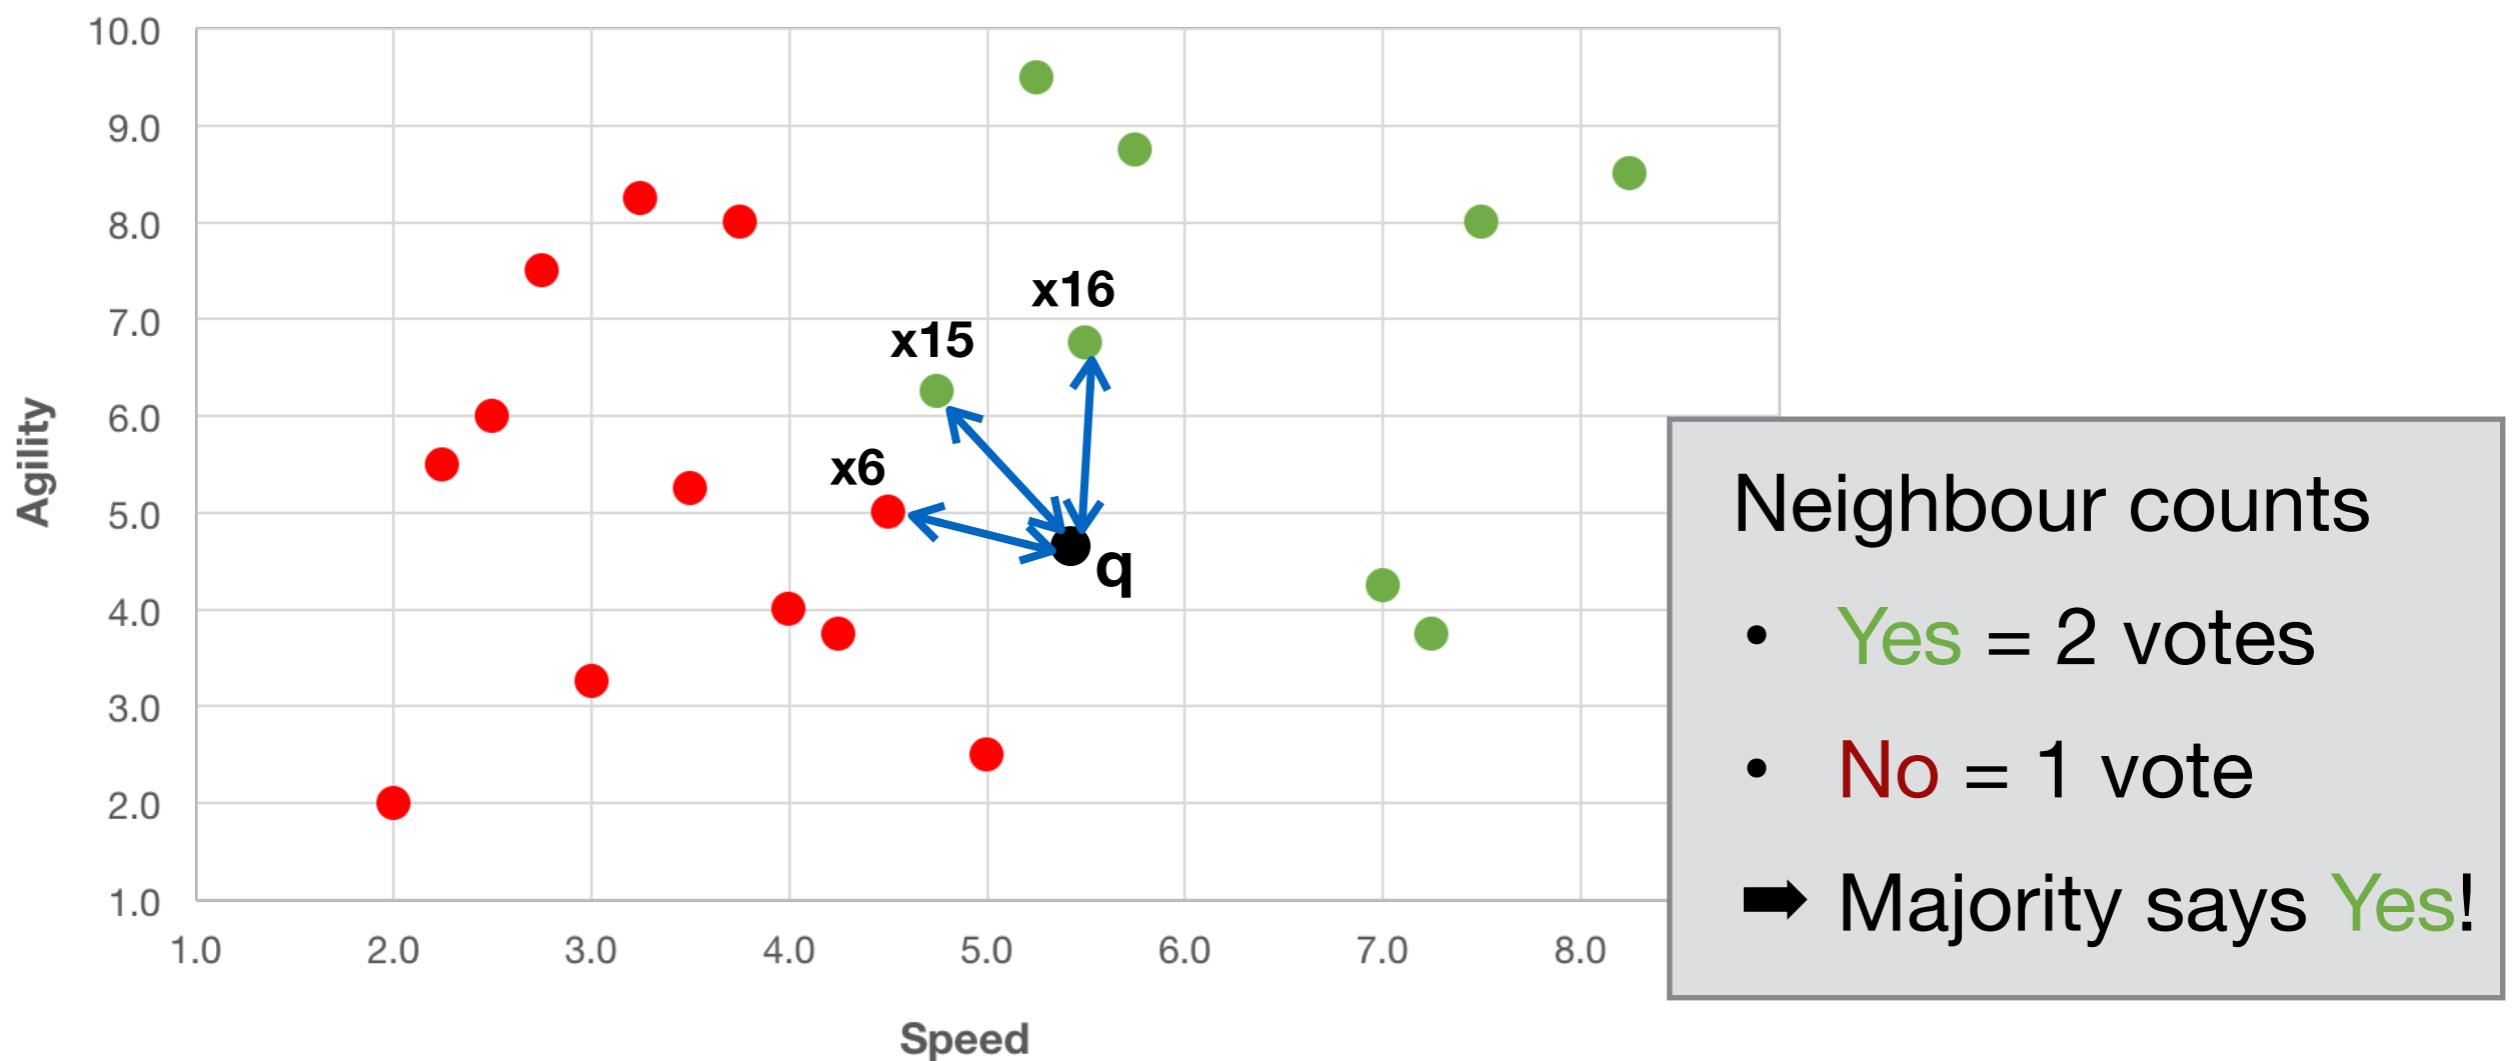
Example: For new query inputs, calculate distance to all training examples. Find $k=3$ nearest examples (i.e. with smallest distances).



k -Nearest Neighbour Classifier

Majority voting: The decision on a label for a new query example is decided based on the “votes” of its k nearest neighbours. The label for the query is the majority label of its neighbours.

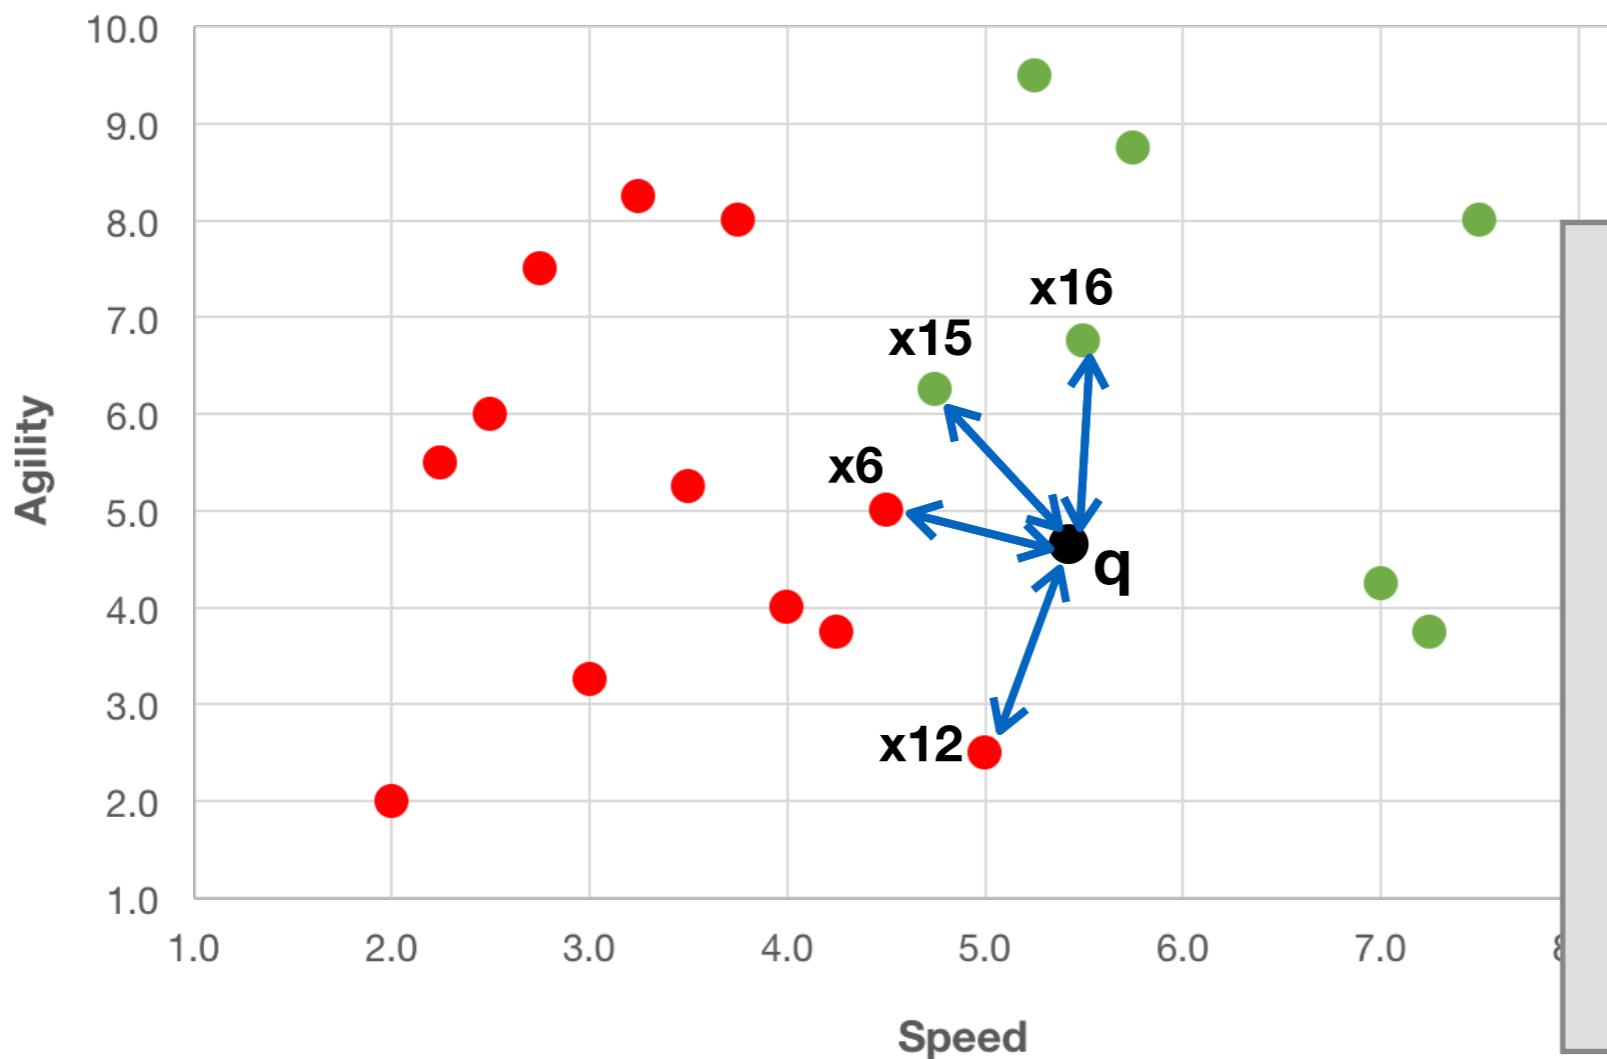
Example: Measure distance from q to all training examples. Find the $k=3$ nearest examples, and use their labels as votes.



k -Nearest Neighbour Classifier

Majority voting: The decision on a label for a new query example is decided based on the “votes” of its k nearest neighbours. The label for the query is the majority label of its neighbours.

Example: Measure distance from q to all training examples. Find the $k=4$ nearest examples, and use their labels as votes.



In the case that...

- Yes = 2 votes
- No = 2 votes

Can break ties...

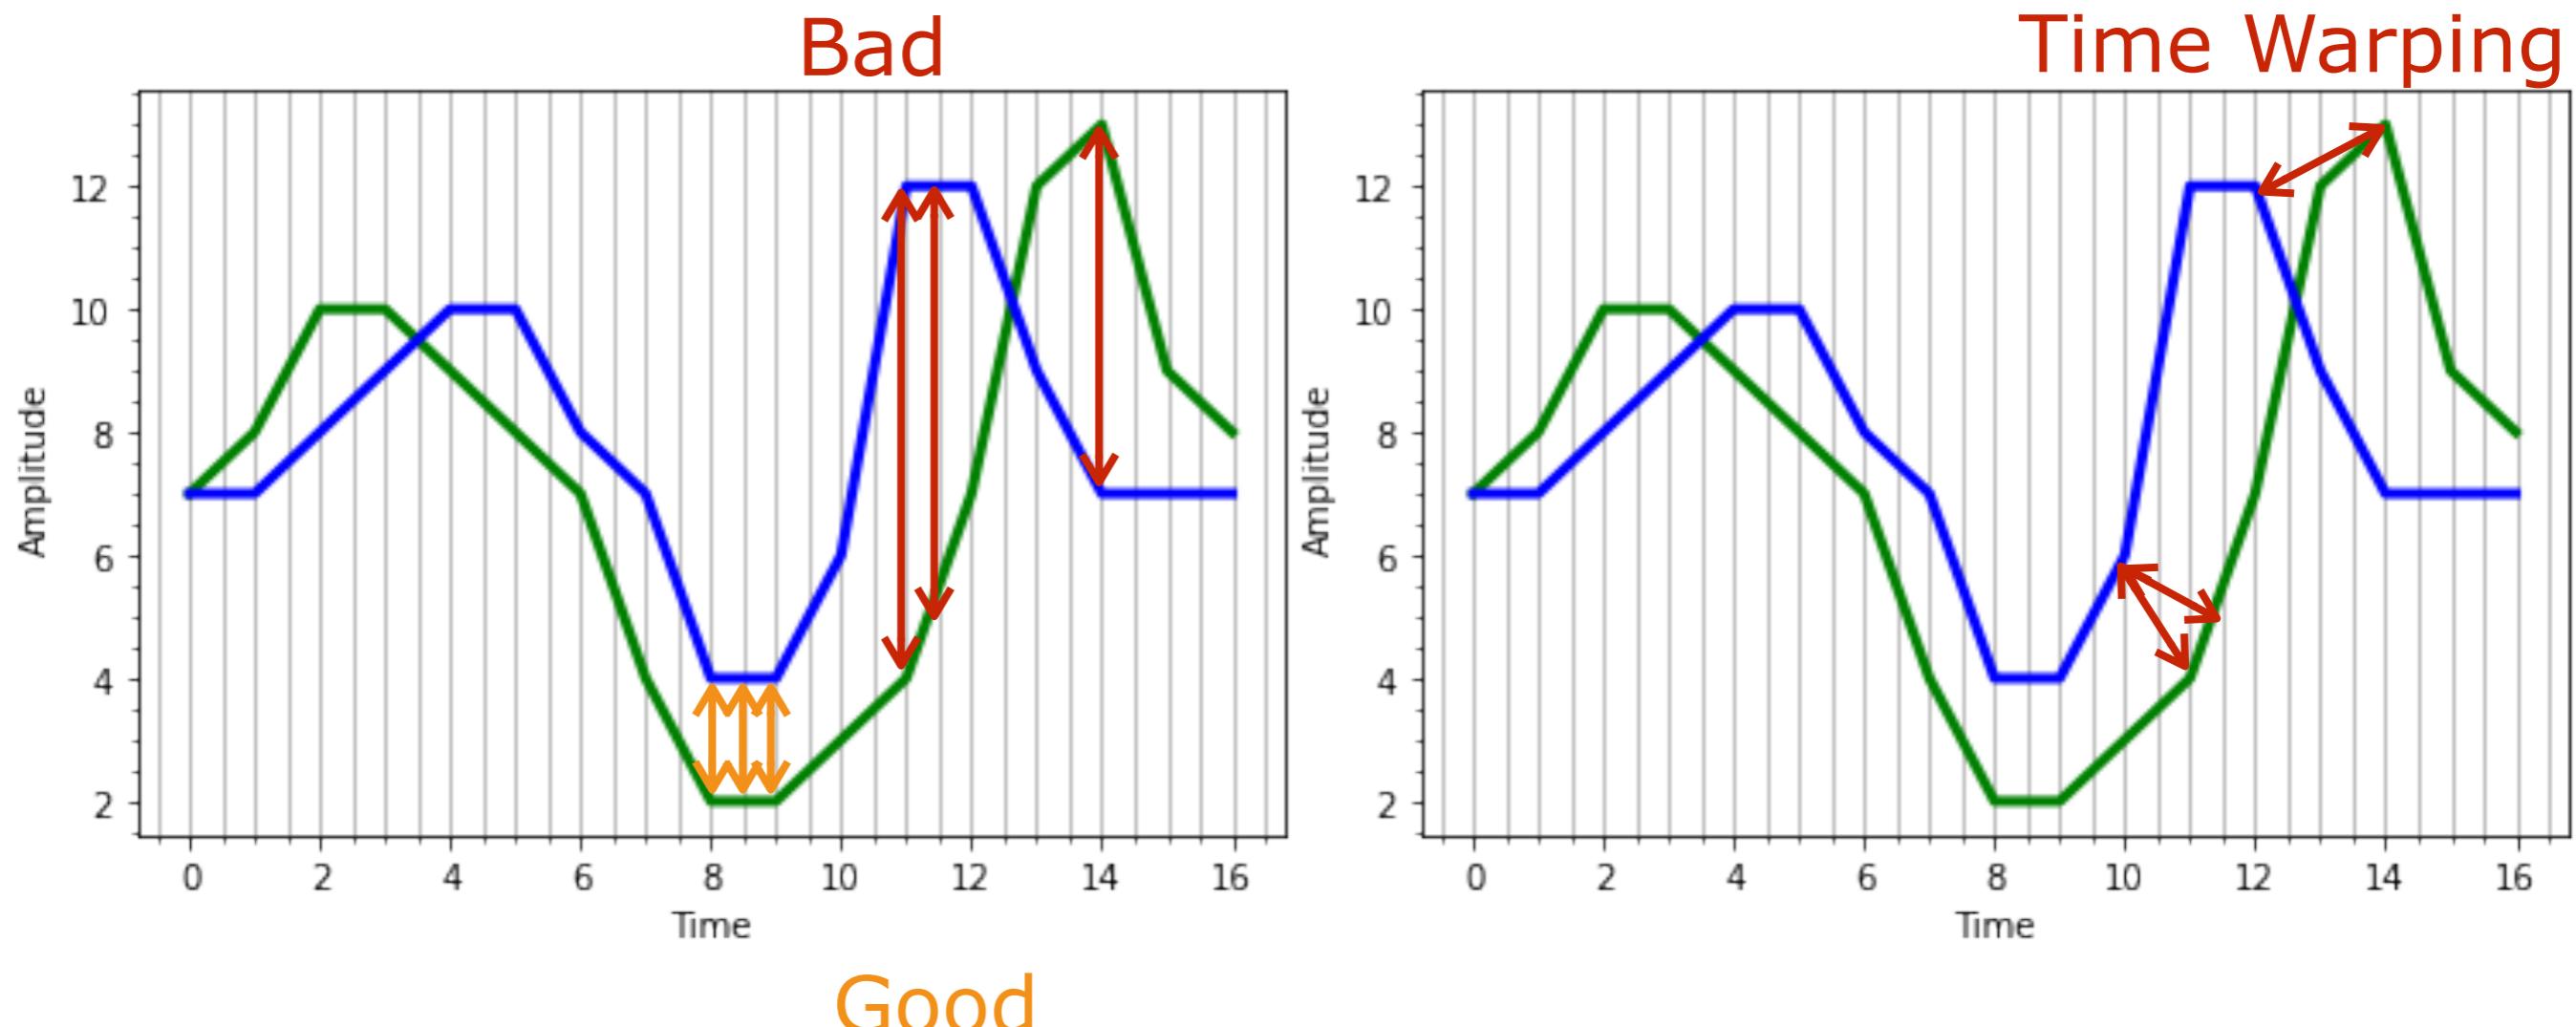
- ▶ At random
- ▶ Based on sum of neighbour distances

Time-Series Classification

- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection

Time-Series as Feature Vectors

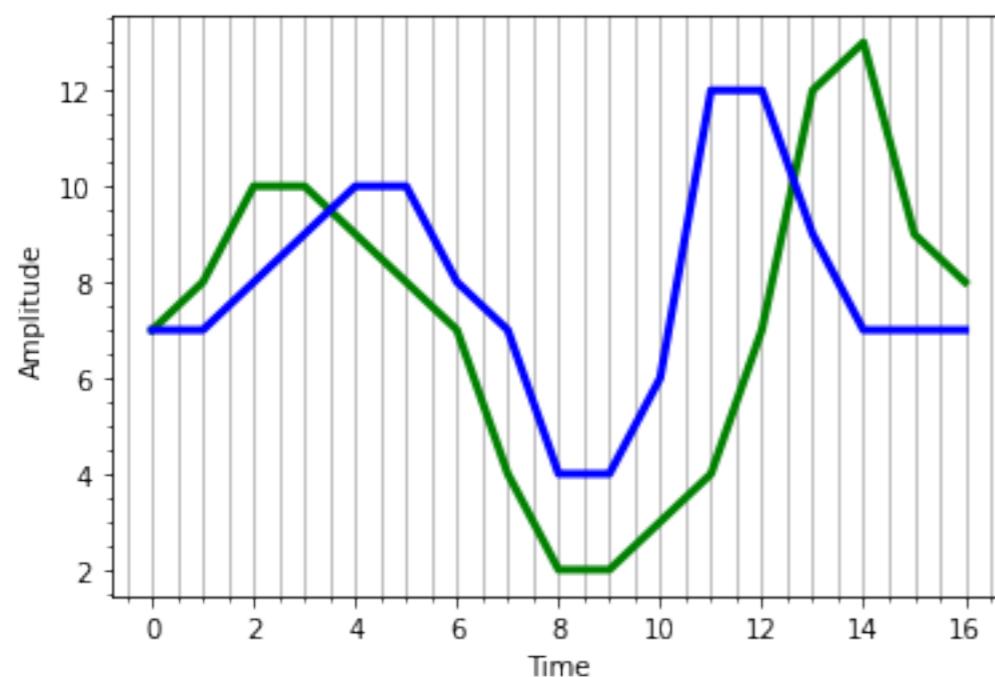
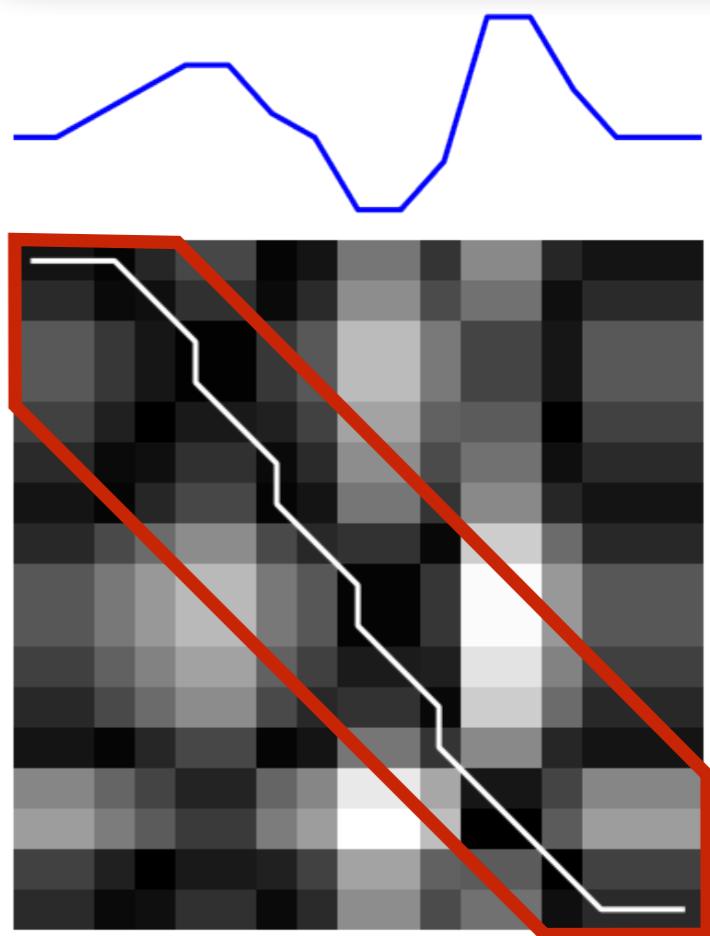
- The problem with Euclidean distance:
 - Time-series length 17 \Rightarrow 17D space



Dynamic Time Warping

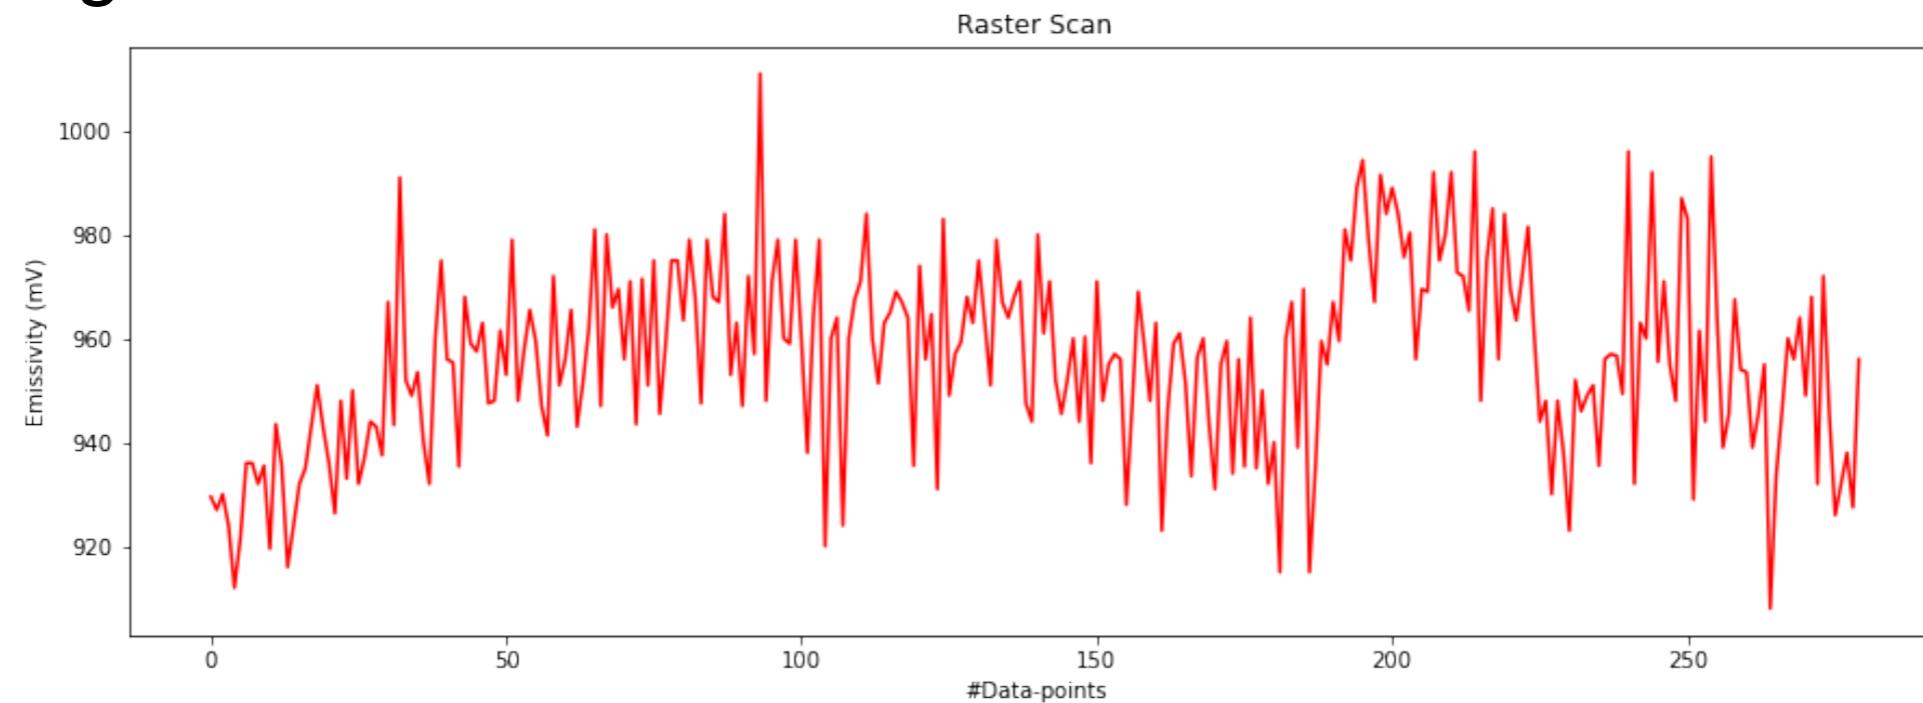
- Allow each point to ‘map’ to any point in the other series
 - Or within a window
- $n \times n$ matrix - find optimum mapping
- $O(n^2)$
- Euclidean distance is $O(n)$

Aside:
 k -NN (and perhaps SVMs) will work with any data format if a similarity/distance measure can be defined.

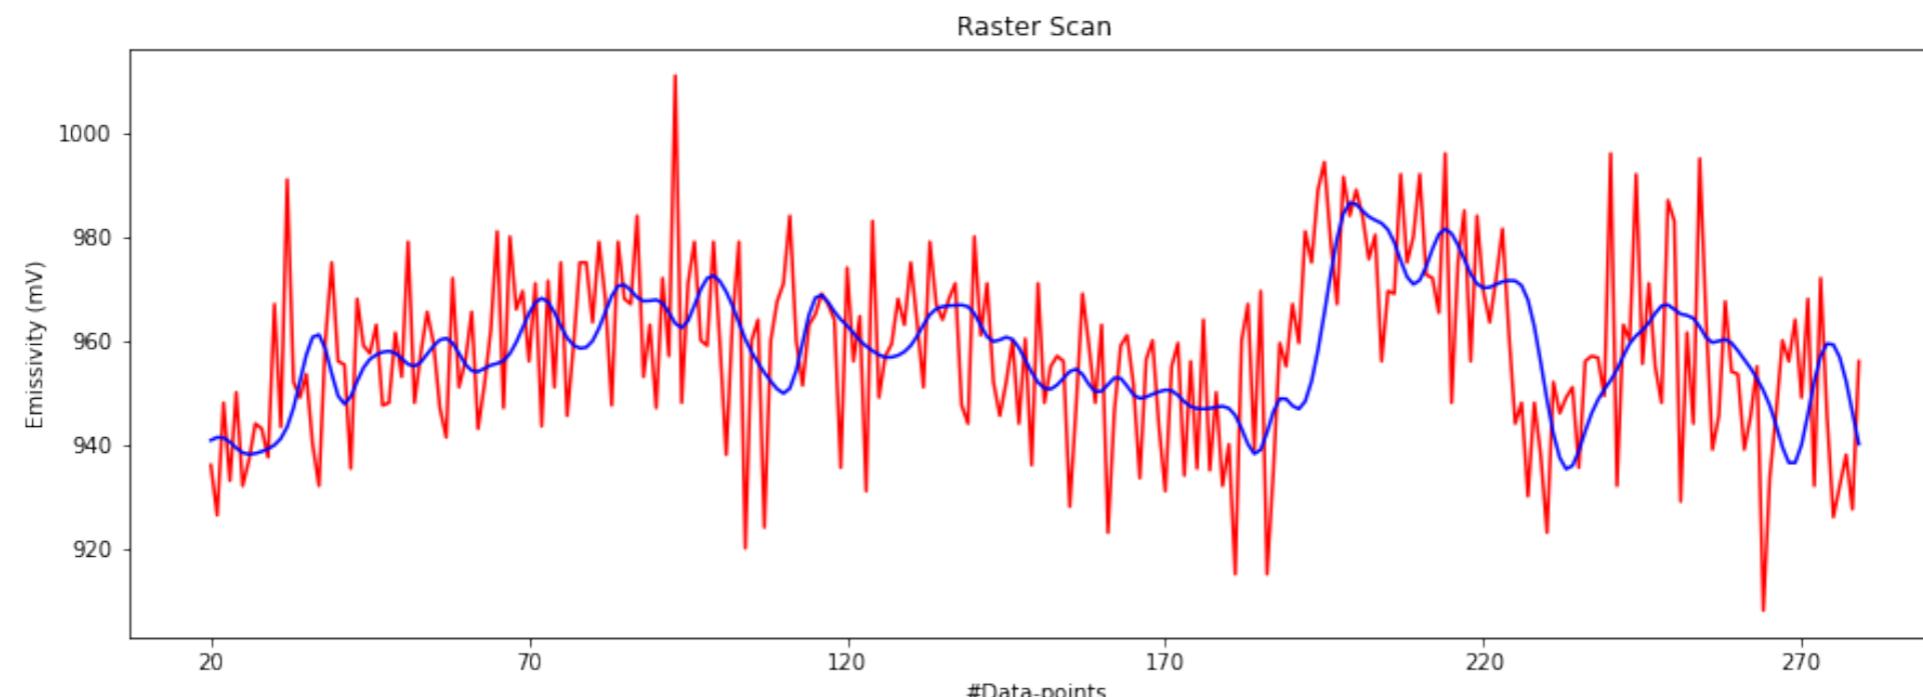


Cleaning the Signal

■ Raw Signal



■ Butterworth Filter



Time-Series Classification

- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection

Performance Measures

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}$$

$$= \frac{4 + 3}{10} = 0.7$$

$$\text{TPRate} = \frac{TP}{TP + FN} = \frac{4}{4 + 1} = 0.8$$

$$\text{FPRate} = \frac{FP}{FP + TN} = \frac{2}{2 + 3} = 0.4$$

$$\text{TNRate} = \frac{TN}{FP + TN} = \frac{3}{2 + 3} = 0.6$$

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{4}{4 + 2} = 0.667$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{4}{4 + 1} = 0.8$$

	Label	Prediction	Correct?	Outcome
1	pore	non-pore		FN
2	pore	pore		TP
3	non-pore	non-pore		TN
4	pore	pore		TP
5	non-pore	pore		FP
6	non-pore	non-pore		TN
7	pore	pore		TP
8	non-pore	pore		FP
9	non-pore	non-pore		TN
10	pore	pore		TP

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$F1 = \frac{2 \times 0.667 \times 0.8}{0.667 + 0.8} = 0.727$$

ROC Analysis

- Varying the decision threshold value θ can lead to different results, and so to different confusion matrices.

	Label	Score	> 0.5?	Prediction	Outcome
1	pore	0.1	N	non-pore	FN
2	pore	0.8	Y	pore	TP
3	non-pore	0.4	N	non-pore	TN
4	pore	0.9	Y	pore	TP
5	non-pore	0.8	Y	pore	FP
6	non-pore	0.2	N	non-pore	TN
7	pore	0.8	Y	pore	TP
8	non-pore	0.6	Y	pore	FP
9	non-pore	0.1	N	non-pore	TN
10	pore	0.9	Y	pore	TP

	Label	Score	> 0.7?	Prediction	Outcome
1	pore	0.1	N	non-pore	FN
2	pore	0.8	Y	pore	TP
3	non-pore	0.4	N	non-pore	TN
4	pore	0.9	Y	pore	TP
5	non-pore	0.8	Y	pore	FP
6	non-pore	0.2	N	non-pore	TN
7	pore	0.8	Y	pore	TP
8	non-pore	0.6	N	non-pore	TN
9	non-pore	0.1	N	non-pore	TN
10	pore	0.9	Y	pore	TP

Decision Threshold $\theta = 0.5$

Predicted Class		
Non	Pore	
TN=3	FP=2	Non
FN=1	TP=4	Pore

Real Class

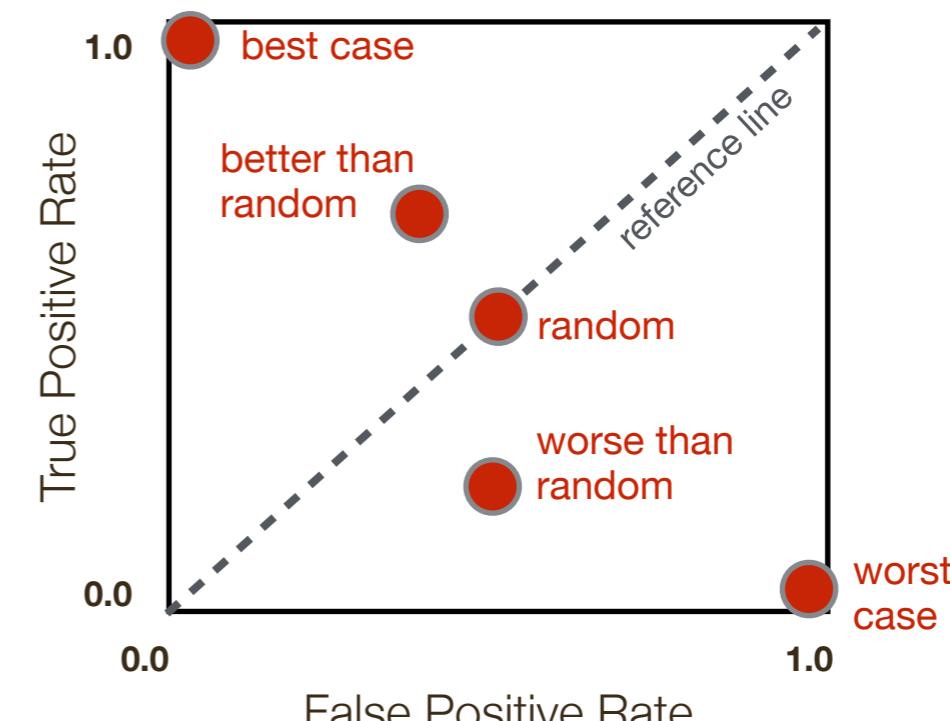
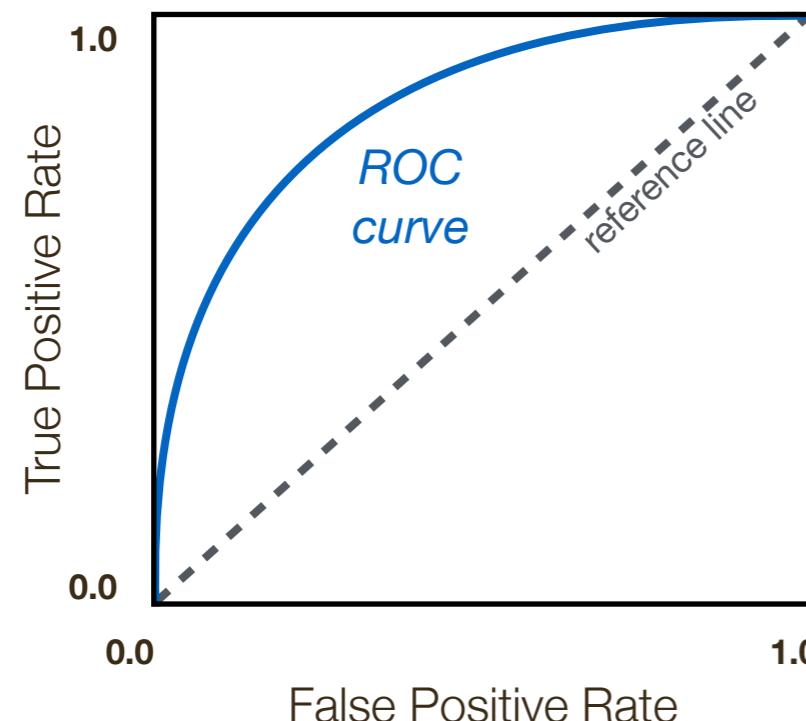
Decision Threshold $\theta = 0.7$

Predicted Class		
Non	Pore	
TN=4	FP=1	Non
FN=1	TP=4	Pore

Real Class

ROC Analysis

- We often want to compare the performance of classifiers at many different decision thresholds (i.e. summarise many confusion matrices).
- A **Receiver Operating Characteristic (ROC Curve)** is a graphical plot of how the true positive rate and false positive rate change over many different thresholds. The curve is drawn by plotting a point for each feasible threshold and joining them.
- A trained classifier should always be above the “random” reference line. The strength of the classifier increases as the ROC curve moves further from the line (i.e. closer to top left corner).



ROC Example

- Given a ranking classifier which will score test samples with a probability P of belonging to the positive class.
- Decision threshold θ controls whether a sample will be classified as positive or negative - i.e. $P > \theta$

Single example $\theta = 0.5$

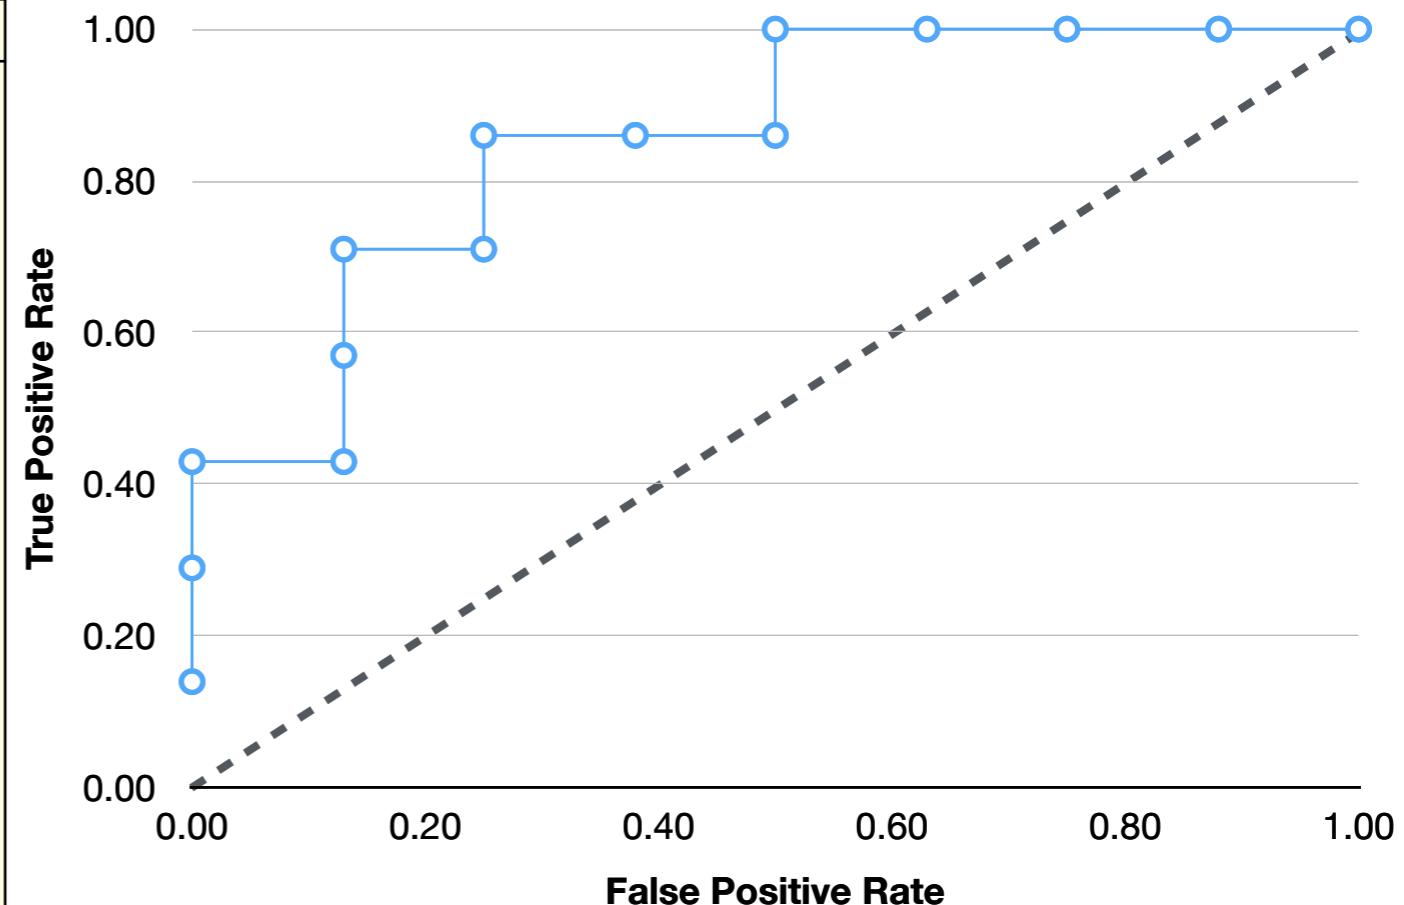
P	> 0.5	Real
0.99	1	1
0.90	1	1
0.85	1	1
0.80	1	0
0.70	1	1
0.70	1	1
0.65	1	0
0.60	1	1
0.45	0	0
0.45	0	0
0.40	0	1
0.30	0	0
0.20	0	0
0.20	0	0
0.20	0	0

ROC Example

- Given a ranking classifier which will score test samples with a probability P of belonging to the positive class.
- Decision threshold θ controls whether a sample will be classified as positive or negative - i.e. $P > \theta$

Single example $\theta = 0.5$

P	> 0.5	Real	FP	TP	FPr	TPr
0.99	1	1	0	1	0.00	0.14
0.90	1	1	0	2	0.00	0.29
0.85	1	1	0	3	0.00	0.43
0.80	1	0	1	3	0.13	0.43
0.70	1	1	1	4	0.13	0.57
0.70	1	1	1	5	0.13	0.71
0.65	1	0	2	5	0.25	0.71
0.60	1	1	2	6	0.25	0.86
0.45	0	0	3	6	0.38	0.86
0.45	0	0	4	6	0.50	0.86
0.40	0	1	4	7	0.50	1.00
0.30	0	0	5	7	0.63	1.00
0.20	0	0	6	7	0.75	1.00
0.20	0	0	7	7	0.88	1.00
0.20	0	0	8	7	1.00	1.00



Time-Series Classification

- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection

Exercise

Check to see if smoothing the data (Butterworth Filter) has actually helped classification accuracy.

Update the ROC curves with results for Euclidean distance and DTW on the unfiltered data.

Hint: Simply replace `X_filtr` with the original `X` data.

Time-Series Classification

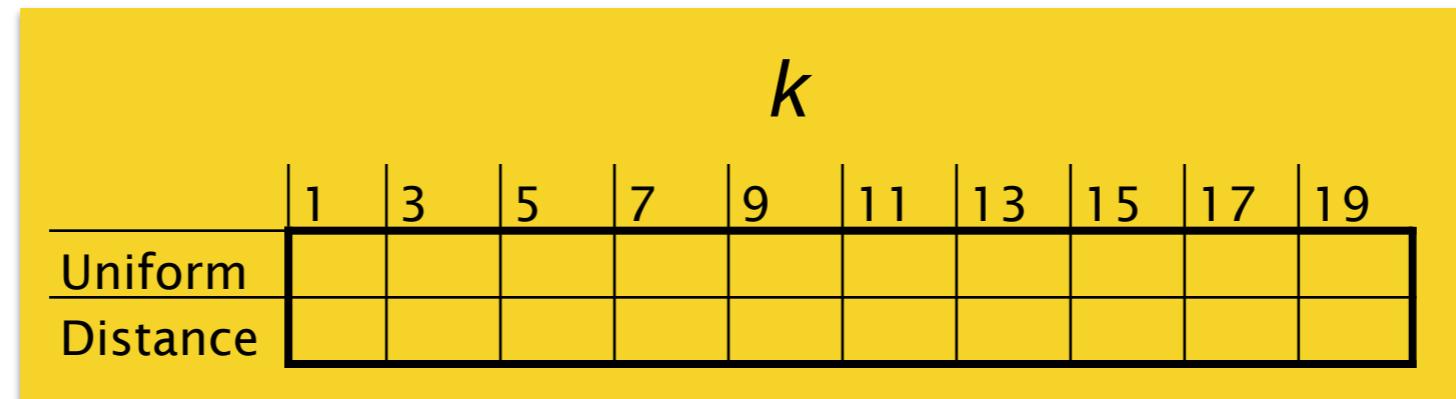
- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection

Model Selection

- AKA: hyper-parameter tuning
- Euclidean distance:
 - k neighbours
 - Distance weighting ?
- DTW:
 - k neighbours
 - Distance weighting
 - Warping constraint:
 - Mechanism, e.g. Sakoe-Chiba
 - Parameters, e.g. Sakoe-Chiba parameters
- Preprocessing
 - Butterworth Filter
 - Data cleaning: missing value imputation, outlier removal

Grid Search		k									
		1	3	5	7	9	11	13	15	17	19
Uniform											
Distance											

Grid Search



- Evaluate alternatives on the training data
- Test data not used in Grid Search (Model Selection)

All Data

Train

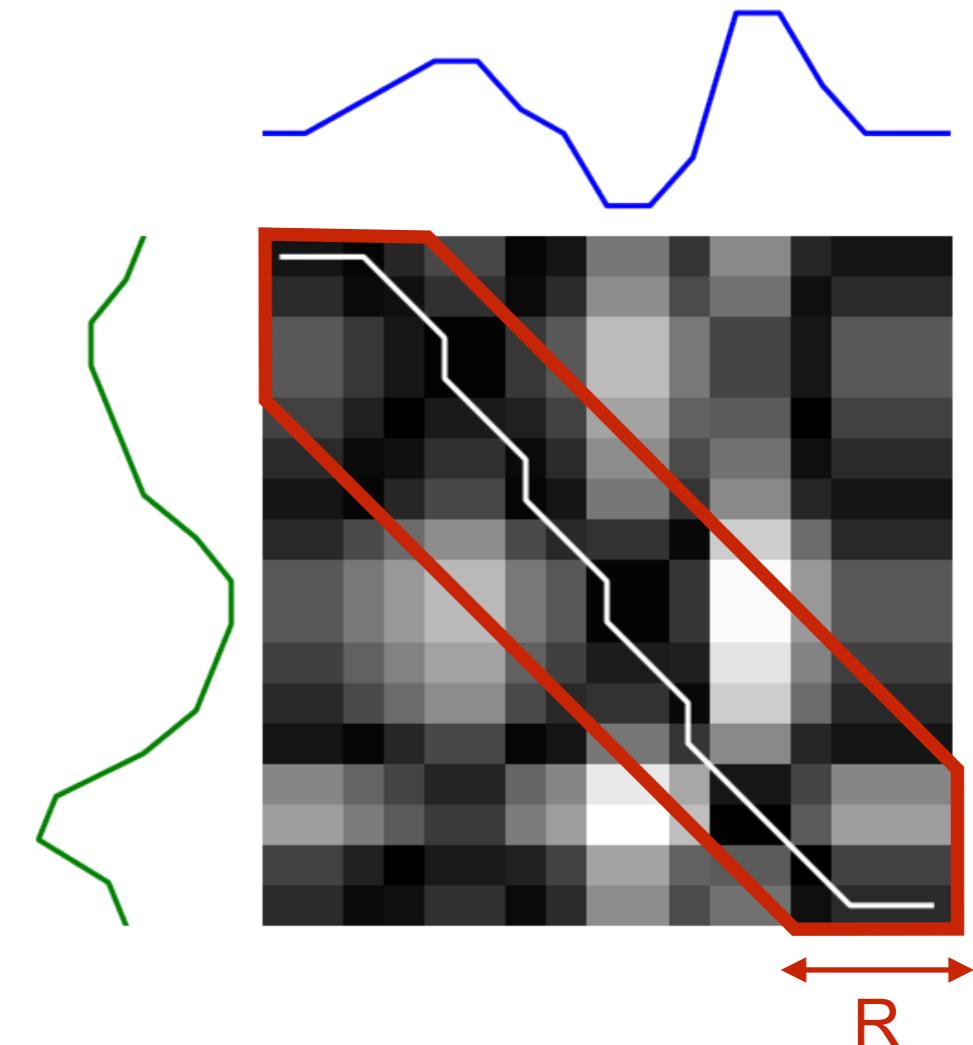
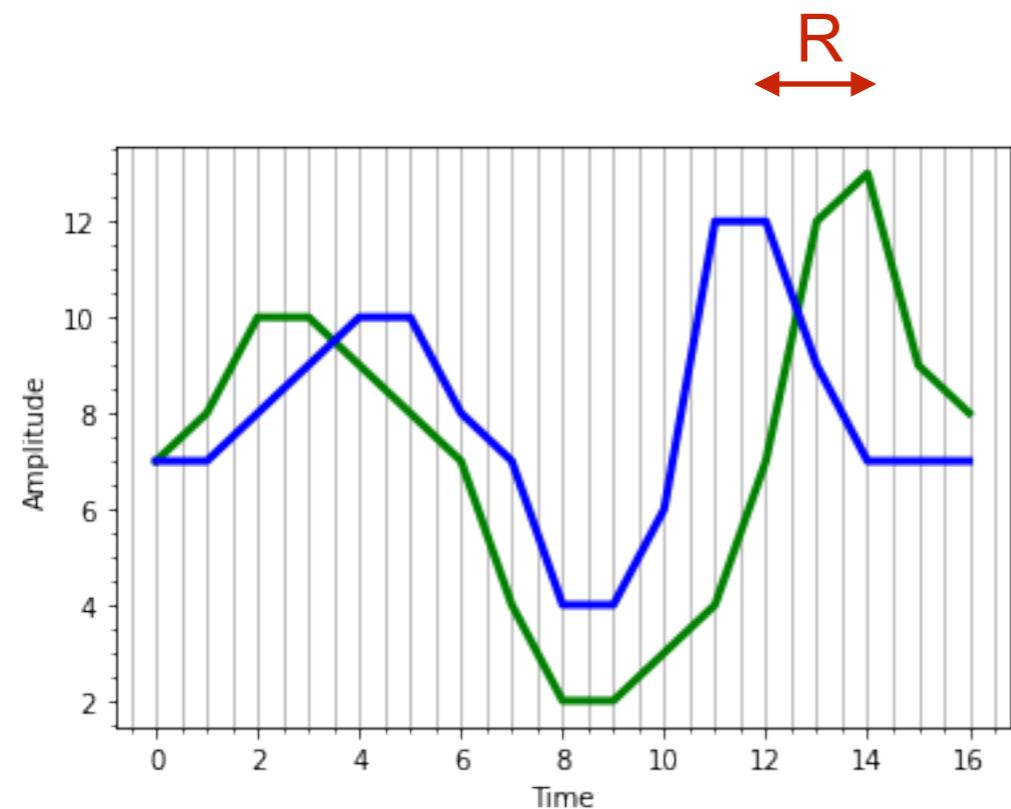
Test

Grid Search
 Cross Validation

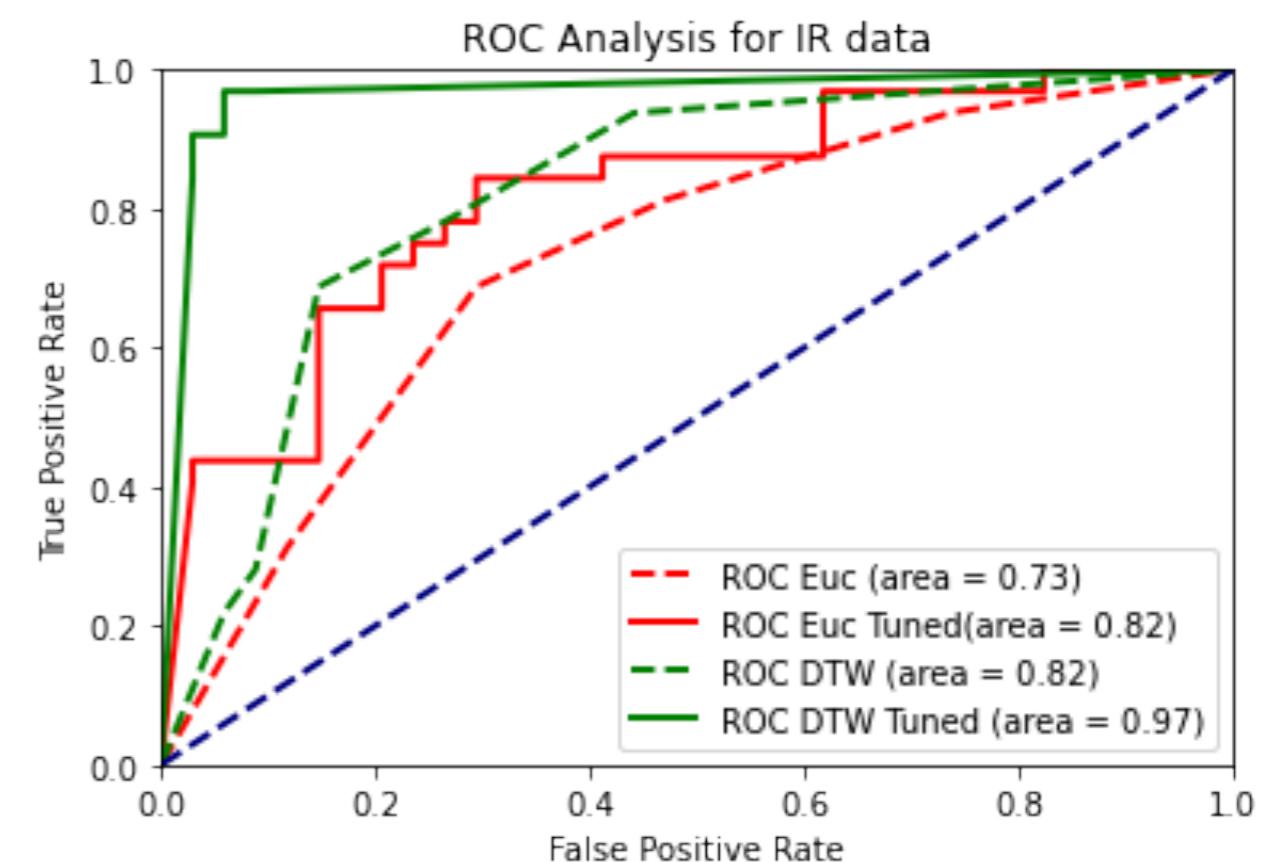
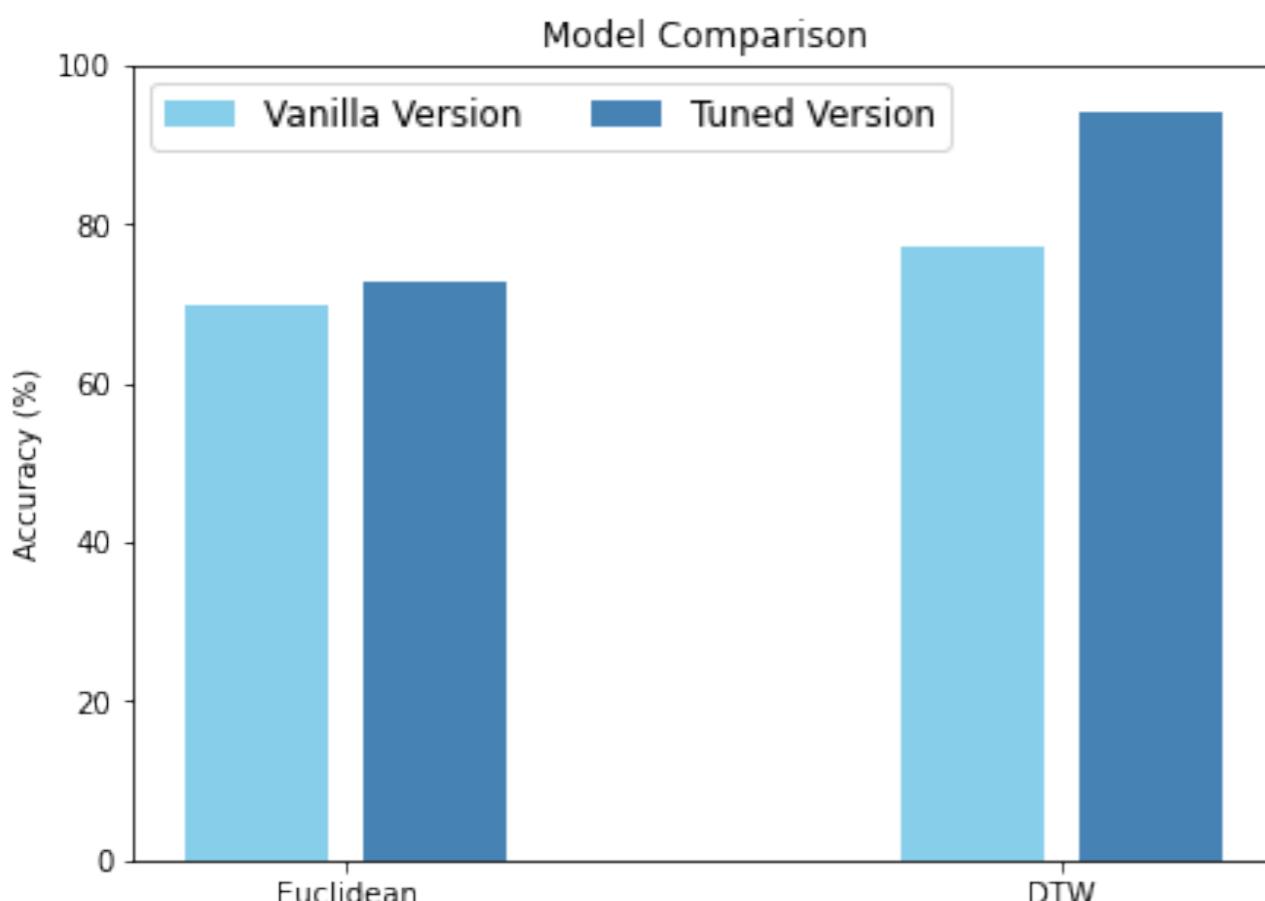
Hold out

DTW Parameters

- Constrain the warping
- Depends on the nature of the data



Final Results



Time-Series Classification

- Manufacturing Context
- Classification in ML
 - Challenges with time-series data
- The k -NN solution
 - Euclidean distance
 - Dynamic Time Warping
- Evaluation
- Exercise
- Model Selection